

Fairness and Proxy Bias in Occupation Prediction: Scaling, Masking, and Mechanistic Attribution across Model Families

Usukhbayar Purevdorj

Yitong Kong

Kirill Kamakaev

Giulio Zhu

Kian Makela

Ambrose Yang

University College London — COMP0087

Abstract

We investigate gender bias and proxy feature encoding in occupation prediction models trained on the Bias-in-Bios dataset. Our study spans two dimensions: *model capacity*, comparing Pythia generative models (160M–1.4B parameters) across zero-shot, few-shot, and fine-tuned regimes, and *model architecture*, contrasting encoder-based classifiers (DistilBERT, RoBERTa) against generative models of comparable performance. We evaluate both task performance (Macro-F1) and gender fairness (Equalized Odds gap, Demographic Parity) under masked and unmasked conditions, and diagnose bias mechanisms through Integrated Gradients attribution validated by top- K token erasure. Our main findings: (1) fine-tuned models substantially outperform zero/few-shot baselines in both performance and fairness; (2) larger models are not inherently fairer—zero-shot bias *increases* with scale; (3) gender-token masking reduces bias for some professions but produces backfire effects for others, suggesting that implicit proxy signals survive removal of explicit gender tokens; (4) explicit gender tokens are not the primary drivers of bias—they appear with weights 61–289 $\times$ smaller than dominant profession tokens, while indirect contextual signals carry the latent gender signal; and (5) proxy bias is model-dependent, the same profession can be stable under one model and unstable under another, though the precise cause (architecture, pre-training, or scale) cannot be isolated from this comparison alone.

1 Introduction

Occupation prediction from biographical text exposes persistent gender stereotypes: models trained on real-world career data reliably associate professions such as *nurse* or *dietitian* with female-typical language patterns, and *surgeon* or *software engineer* with male-typical ones. De-Arteaga et al.

[2019] demonstrated on the Bias-in-Bios dataset that scrubbing explicit gender signals does not eliminate these disparities—implicit proxy features survive.

This paper addresses three open questions: how does model *capacity* interact with fairness; how does model *architecture* shape which proxy features are learned; and what are those features precisely.

Contributions. *Primary novel contribution* (to our knowledge, no prior work on this dataset provides this): (i) token-level Integrated Gradients attribution tables stratified by both profession and gender, with proxy stability compared across encoder models. *Secondary novel contribution*: (ii) an erasure faithfulness validation of those attribution scores on the Bias-in-Bios task. *Empirical applications of known methods to this setting*: (iii) a Pythia scaling study (160M–1.4B) across zero-shot, few-shot, and fine-tuned regimes; (iv) a cross-architecture comparison of encoder and generative model fairness on Bias-in-Bios; (v) a paired masked/unmasked evaluation of gender-token scrubbing on both task performance and fairness metrics.

Key Findings. Critically, zero-shot bias *increases* with Pythia scale (EO gap ~ 0.09 at 160M to ~ 0.21 at 1.4B), gender tokens carry weights 61–289 $\times$ smaller than dominant profession tokens [De-Arteaga et al., 2019], and—novel on this dataset—proxy stability is model-dependent: the same profession can be stable under one encoder and unstable under another. Full findings are reported in §6.2 and §6.3. We do not claim these findings generalise beyond the Bias-in-Bios setting.

2 Related Work

Bias in occupation prediction. De-Arteaga et al. [2019] introduced Bias-in-Bios and showed that classifiers exhibit gender disparities even after pro-

noun scrubbing. Romanov et al. [2019] analysed limits of de-biasing through scrubbing. We extend these findings by identifying and quantifying specific proxy tokens through gradient-based attribution.

Fairness metrics. Hardt et al. [2016] formalised Equalized Odds; Dwork et al. [2012] introduced demographic parity. We evaluate both using a one-vs-rest multiclass extension macro-averaged over occupations.

Attribution methods. Sundararajan et al. [2017] introduced Integrated Gradients satisfying completeness and sensitivity axioms. Attribution methods have been applied to gender bias in coreference [Attanasio et al., 2022] but their systematic use in occupation classification is limited. Critically, no prior work on the Bias-in-Bios occupation prediction task provides token-level proxy word tables with quantified attribution weights, nor examines how those weights shift under gender masking.

Model scaling and fairness. Bommasani et al. [2021] discusses risks of foundation models including bias amplification. Our scaling study provides empirical evidence that zero-shot bias increases with model size in the Pythia family.

Gender masking. Lu et al. [2020] study counterfactual data augmentation for gender bias. Unlike prior work, we measure not only the effect of masking on fairness metrics but also on token-level attribution patterns — providing the first internal view of how occupation prediction models redistribute feature importance when gender signals are removed.

3 Dataset and Preprocessing

3.1 Bias-in-Bios

We use the Bias-in-Bios dataset [De-Arteaga et al., 2019] (LabHC/bias_in_bios), a standard NLP fairness benchmark. Each example is a biography (`hard_text`) with an integer profession label and a binary gender attribute (inferred from pronoun usage by the original authors; $0 \rightarrow \text{M}$, $1 \rightarrow \text{F}$). We restrict to the **top 20 most frequent occupations** from the training split. The task is to predict occupation from biography text alone; gender is a post-hoc evaluation attribute only.

Table 2 shows the per-profession counts and gender distribution in the training split. Several occupations exhibit strong gender skew: *dietitian* (92.9% F), *nurse* (90.8% F), and *model* (82.7% F) are heavily

Split	Total	Female	Male
Train	248,141	46.16%	53.84%
Dev	38,198	46.16%	53.84%
Test	95,468	46.17%	53.83%
Total	381,807	46.16%	53.84%

Table 1: Dataset split sizes and gender distribution. The near-identical F% across splits (range: 0.01%, i.e. approximately 46.2% female in each) indicates the original dataset was stratified by gender when partitioned.

ily female-dominated, while *surgeon* (14.8% F), *software engineer* (15.8% F), and *composer* (16.4% F) are heavily male-dominated.

Profession	N (train)	F%
professor	76,748	45.1
physician	26,648	49.4
attorney	21,169	38.3
photographer	15,773	35.7
journalist	12,960	49.5
nurse	12,316	90.8
psychologist	11,945	62.1
teacher	10,531	60.2
dentist	9,479	35.3
surgeon	8,829	14.8
architect	6,568	23.7
painter	5,025	45.7
model	4,867	82.7
filmmaker	4,545	32.9
poet	4,558	49.0
software engineer	4,492	15.8
composer	3,637	16.4
accountant	3,660	36.7
dietitian	2,567	92.9
comedian	1,824	21.1

Table 2: Per-profession training counts and F% ($M\% = 100 - F\%$). Distribution is consistent across all splits.

3.2 Evaluation Subset

All models are evaluated on a canonical **3,000-sample test subset**, stratified proportionally by profession from the 95,468-example test split; a shared ID list ensures cross-model comparability. Per-profession sizes range from 932 (professor) to 27 (comedian); results for low-frequency professions should be interpreted with caution.

3.3 Gender Masking

Two conditions: **Unmasked** (original text) and **Masked** (gender tokens replaced with the model’s

native mask token—[MASK] for DistilBERT, <mask> for RoBERTa—via case-insensitive word-boundary regex). Three token categories are masked: pronouns and possessives (*he, she, his, her, etc.*), titles and honorifics (*mr, mrs, ms, sir, lady, etc.*), and 35 gendered nouns (*father, mother, husband, wife, actor, actress, etc.*). Name masking was not applied: general name rules introduce uncontrolled false positives, and names carry biographical content beyond gender that would confound gender-signal removal with loss of biographical information.

4 Methods

4.1 Pythia: Zero-shot and Few-shot

We evaluate Pythia [Biderman et al., 2023] at 160M, 410M, and 1.4B parameters under zero-shot (biography + candidate list) and few-shot ($K = 4$ exemplar pairs) regimes, with and without gender masking. As a causal language model, Pythia scores each candidate label c_k by its joint log-probability given prompt p :

$$\log P(c_k | p) = \sum_{t=1}^{|c_k|} \log P(c_k^{(t)} | p, c_k^{(1)}, \dots, c_k^{(t-1)}) \quad (1)$$

The predicted label is $\hat{y} = \arg \max_{c_k} \log P(c_k | p)$.

4.2 Pythia: Supervised Fine-tuning

Pythia is fine-tuned as a sequence classifier (linear head over the final hidden state) using LoRA [Hu et al., 2022] on attention projections ($r = 16$, $\alpha = 32$, dropout 0.1), lr 2×10^{-5} , batch 16, 3 epochs, max length 256; best checkpoint selected by dev Macro-F1. QLoRA (4-bit) is used for 1.4B.

4.3 Encoder Baselines: DistilBERT and RoBERTa

DistilBERT [Sanh et al., 2019] (66M) and RoBERTa-base [Liu et al., 2019] (125M) are fine-tuned as sequence classifiers with a linear head, lr 2×10^{-5} , batch 16, 3 epochs, max length 256, best checkpoint by Macro-F1. The same pipeline runs for both masked and unmasked conditions with a shared label mapping.

4.4 Evaluation Metrics

We evaluate models on both task performance and fairness. Let Y represent the true occupation label, $\hat{Y}$ the predicted label, $A \in \{M, F\}$ the binary

gender attribute, and $\mathcal{C}$ the set of 20 occupation classes.

Task Performance: We report overall Accuracy and Macro-F1 score to account for the heavy class imbalance present in the dataset.

Fairness Metrics: We measure gender disparities using one-vs-rest multiclass extensions of Demographic Parity (DP) and Equalized Odds (EO), macro-averaged across all occupations.

Demographic Parity requires the prediction to be independent of the sensitive attribute. We compute the macro-averaged DP gap as the absolute difference in positive prediction rates between genders for each class:

$$\Delta DP = \frac{1}{|\mathcal{C}|} \sum_{c \in \mathcal{C}} \left| P(\hat{Y} = c | A = M) - P(\hat{Y} = c | A = F) \right|$$

Equalized Odds requires the prediction to be conditionally independent of the sensitive attribute given the true label. To compute the EO gap, we first calculate the macro-averaged True Positive Rate (TPR) gap and False Positive Rate (FPR) gap:

$$\Delta TPR = \frac{1}{|\mathcal{C}|} \sum_{c \in \mathcal{C}} \left| P(\hat{Y} = c | Y = c, A = M) - P(\hat{Y} = c | Y = c, A = F) \right|$$

$$\Delta FPR = \frac{1}{|\mathcal{C}|} \sum_{c \in \mathcal{C}} \left| P(\hat{Y} = c | Y \neq c, A = M) - P(\hat{Y} = c | Y \neq c, A = F) \right|$$

The final Equalized Odds gap is defined as the maximum of these two macro-averaged disparities:

$$EO \text{ Gap} = \max(\Delta TPR, \Delta FPR)$$

4.5 Proxy Bias Attribution

Integrated Gradients [Sundararajan et al., 2017] is applied to the two fine-tuned encoder models, with attribution computed toward the *predicted* class index. Attributions are computed using **32 interpolation steps** and a **zero-embedding baseline**. Token-level scores are obtained by **summing** attribution values across the embedding dimension for each token position. Subword tokens are merged to word level before top- K selection: BERT-style ## continuation pieces and whitespace-prefixed tokens (RoBERTa’s byte-pair encoding marks word

boundaries with a special prefix character) are handled separately, with scores summed across pieces. Only tokens with **positive** attribution scores and minimum length of 2 characters are retained. Maximum sequence length is **256 tokens**.

The pipeline operates in two stages with distinct K values. **Stage 1 — per-example extraction** ($K_1=5$): the 5 highest-scoring positive tokens are retained per example. **Stage 2 — profession-level aggregation**: tokens appearing across all examples for a profession are aggregated (minimum length 3, stopwords removed) to compute:

$$\text{wscore}(t, p) = \text{count_topk}(t, p) \times \overline{\text{attr}}(t, p) \quad (2)$$

where $\text{count_topk}(t, p)$ is the number of examples in which token t appears in the stage-1 top-5 for profession p , and $\overline{\text{attr}}(t, p)$ is the mean attribution score. Aggregation uses the model’s *predicted* label (not the true label), consistent with attribution targeting the predicted class.

Stability analysis. The masked vs. unmasked comparison takes the **top-10 aggregated tokens** per profession ($K_2=10$) and computes the shared token ratio:

$$r(p) = \frac{|\mathcal{T}_{\text{sh}}|}{|\mathcal{T}_{\text{sh}}| + |\mathcal{T}_{\text{m}}| + |\mathcal{T}_{\text{u}}|} \quad (3)$$

where $\mathcal{T}_{\text{sh}}$, $\mathcal{T}_{\text{m}}$, $\mathcal{T}_{\text{u}}$ are the sets of tokens shared, masked-only, and unmasked-only respectively among the top-10. Gender-conditioned proxy tables are produced identically, stratified by profession and gender field. **Faithfulness validation.** The top-5 attributed tokens are erased from each example’s text and the model is re-run to measure the probability drop:

$$\text{Faith}(x) = P(y | x) - P(y | x_{\text{erased}}) \quad (4)$$

5 Experimental Setup

Code and prediction files are available at <https://github.com/usukhbayarp/comp0087-fairness-occupation-prediction>.

Table 3 summarises all evaluated configurations.

6 Results and Discussion

We evaluate our models across three dimensions: task performance (§6.1), gender fairness (§6.2), and token-level proxy bias mechanisms (§6.3).

The following results are descriptive of only the 3,000-sample evaluation subset, which may introduce some noise.

Model	Type	Regime	Cond.
Pythia-160M	Gen.	zero-shot, few-shot, FT	M, U
Pythia-410M	Gen.	zero-shot, few-shot, FT	M, U
Pythia-1.4B	Gen.	zero-shot, few-shot, FT	M, U
DistilBERT	Enc.	Fine-tuned	M, U
RoBERTa	Enc.	Fine-tuned	M, U

Table 3: All model configurations. FT=fine-tuned; M=masked; U=unmasked.

Model	C.	Acc.	F1
Pythia-160M FT	U	0.079	0.020
Pythia-160M FT	M	0.119	0.023
Pythia-410M FT	U	0.581	0.323
Pythia-410M FT	M	0.576	0.325
Pythia-1.4B FT	U	0.877	0.848
Pythia-1.4B FT	M	0.877	0.852
DistilBERT FT	U	0.875	0.843
DistilBERT FT	M	0.872	0.844
RoBERTa FT	U	0.882	0.851
RoBERTa FT	M	0.880	0.851
Pythia-160M zero-shot	U	0.405	0.230
Pythia-160M zero-shot	M	0.400	0.223
Pythia-160M few-shot	U	0.241	0.202
Pythia-160M few-shot	M	0.249	0.206
Pythia-410M zero-shot	U	0.536	0.557
Pythia-410M zero-shot	M	0.533	0.553
Pythia-410M few-shot	U	0.366	0.391
Pythia-410M few-shot	M	0.363	0.389
Pythia-1.4B zero-shot	U	0.658	0.567
Pythia-1.4B zero-shot	M	0.659	0.568
Pythia-1.4B few-shot	U	0.558	0.578
Pythia-1.4B few-shot	M	0.558	0.577

Table 4: Task performance results. U=unmasked; M=masked.

6.1 Task Performance

Pythia-1.4B fine-tuned (M) reaches F1=0.877, nearing encoder baselines (RoBERTa (M) 0.880, DistilBERT (M) 0.872); Pythia-160M fine-tuned fails (F1 $\approx$ 0.02), likely due to label-space collapse under LoRA at insufficient capacity. Zero-shot F1 scales from 0.23 (160M) to 0.57 (1.4B); few-shot hurts at smaller scales (410M $\rightarrow$ 160M: 0.557 $\rightarrow$ 0.391) but recovers at 1.4B (0.578), suggesting a minimum capacity threshold. Masking has negligible impact on task performance across all models.

6.2 Fairness Results

Fine-tuned models achieve markedly lower EO gaps than zero-shot counterparts; zero-shot EO grows from $\sim$ 0.09 (160M) to $\sim$ 0.21 (1.4B), am-

Model	C.	EO	DP	TPR	FPR
RoBERTa FT	U	0.082	0.024	0.082	0.005
RoBERTa FT	M	0.063	0.022	0.063	0.004
DistilBERT FT	U	0.088	0.022	0.088	0.004
DistilBERT FT	M	0.071	0.022	0.071	0.004
Pythia-1.4B FT	U	0.086	0.023	0.086	0.004
Pythia-1.4B FT	M	0.080	0.022	0.080	0.003
Pythia-1.4B FS	U	0.121	0.031	0.121	0.016
Pythia-1.4B FS	M	0.124	0.031	0.124	0.016
Pythia-1.4B 0s	U	0.210	0.029	0.210	0.019
Pythia-1.4B 0s	M	0.210	0.029	0.210	0.019
Pythia-410m FT	U	0.085	0.012	0.085	0.006
Pythia-410m FT	M	0.088	0.011	0.088	0.006
Pythia-410m FS	U	0.126	0.021	0.126	0.012
Pythia-410m FS	M	0.128	0.020	0.128	0.011
Pythia-410m 0s	U	0.143	0.037	0.143	0.020
Pythia-410m 0s	M	0.131	0.037	0.131	0.021
Pythia-160m FT	U	0.017	0.001	0.017	0.001
Pythia-160m FT	M	0.003	0.002	0.003	0.002
Pythia-160m FS	U	0.080	0.012	0.080	0.007
Pythia-160m FS	M	0.077	0.013	0.077	0.008
Pythia-160m 0s	U	0.092	0.027	0.092	0.025
Pythia-160m 0s	M	0.092	0.027	0.092	0.024

Table 5: Fairness metrics (EO = Equalized Odds gap, DP = Demographic Parity gap, TPR = True Positive Rate gap, FPR = False Positive Rate gap) for each model, both masked (M) and unmasked (U). Row label abbreviations: FT = fine-tuned; FS = few-shot; 0s = zero-shot.




Figure 1: Pareto frontier of Macro-F1 vs. Equalized Odds gap. RoBERTa masked fine-tuned, followed by DistilBERT, is the most Pareto-optimal model. The generative models which are masked and with largest capacity, as well as the baseline encoders, strongly dominate lower-capacity, zero-shot and few-shot alternatives.

plifying stereotypes without supervision. Masking reduces aggregate EO for encoders (RoBERTa: 0.082→0.063; DistilBERT: 0.088→0.071) but is non-uniform at profession level: Pythia-1.4B FT sees backfire for *teacher*, *poet*, and *painter*,




Figure 2: Scaling behavior of Equalized Odds gap across model sizes, 1.4B (red), 410m (green), and 160m (blue). The solid line corresponds to unmasked and dashed to masked. Zero-shot bias increases with scale, while fine-tuning mitigates this effect.

while Pythia-410M FT shows aggregate backfire (0.085→0.088). TPR gap equals EO gap across all models, indicating that differential true positive rates are the primary fairness issue. To understand why disparities persist under masking, we turn to token-level attribution analysis.

6.3 Proxy Bias Analysis

We analyse which tokens drive occupation predictions and how their importance shifts under gender masking. We first assess attribution stability under masking, then examine which token types carry the gender signal, and finally assess token importance using erasure-based faithfulness analysis.

Profession	DistilBERT	RoBERTa
comedian	0.54	0.33
dentist	0.54	0.33
dietitian	0.67	0.54
nurse	0.43	0.54
psychologist	0.82	0.82
software engineer	0.67	0.43
attorney	0.67	0.82
journalist	0.67	0.67
physician	0.67	0.54
professor	0.82	1.00

Table 6: Shared token ratio (masked vs. unmasked). Higher = more stable.

Under RoBERTa, *comedian* and *dentist* are the most unstable professions (ratio 0.33), while under DistilBERT both are moderately stable (0.54)—the same reversal observed for *nurse* (DistilBERT 0.43,

RoBERTa 0.54). *Professor* reaches maximum stability under RoBERTa (1.00), meaning its top-10 tokens are identical across masked and unmasked conditions. These cross-profession reversals reinforce the conclusion that proxy stability depends on the specific model, not only training data.

When gender tokens are removed, attribution is redistributed across tokens rather than concentrated on the same dominant token, consistent with the completeness property of Integrated Gradients. The profession label token (e.g. *nurse* for the nurse class) frequently appears among the top- K attributed tokens, reflecting label leakage where the profession name occurs directly in the biography. This is expected; the methodologically informative signal lies in how attribution shifts between unmasked and masked conditions. For DistilBERT on *nurse*, the label token *nurse* remains prominent under both conditions (129.89 unmasked, 152.73 masked), consistent with label leakage noted above. The informative shift occurs in secondary tokens: *practitioner* declines from 26.51 to 11.18 and *graduated* from 19.11 to 1.47, while *hospital* (8.96) and *especially* (15.52) gain attribution weight. For RoBERTa on *physician*, masking shifts weight toward higher-level clinical context: *medicine* rises from 166.64 to 292.25 and *patients* (52.72) emerges as a new signal, while *affiliated* (68.78 unmasked) drops from the top-5.

Model	Token	Type	W.Score
DistilBERT	<i>actor</i>	gender	1.15
DistilBERT	<i>sister</i>	gender	1.13
DistilBERT	<i>mrs</i>	gender	1.05
RoBERTa	<i>mrs</i>	gender	2.11
RoBERTa	<i>wife</i>	gender	0.78
RoBERTa	<i>nursing</i>	profession	70.44
RoBERTa	<i>medicine</i>	profession	166.64
RoBERTa	<i>research</i>	profession	333.32

Table 7: Attribution weight of gender tokens vs. profession tokens. Gender tokens appear in both models but with substantially smaller weights than dominant profession tokens.

Under both models, a small number of explicit gender tokens appear in the aggregated top- K features, but with substantially smaller weights than dominant profession tokens: DistilBERT’s highest gender token (*actor*: 1.15) is $61\text{--}289\times$ smaller than the profession tokens in Table 7 (70.44–333.32), and RoBERTa’s *mrs* (2.11) is similarly dwarfed.

Bias likely operates primarily through indirect contextual signals rather than explicit gender indicators.

Gender-conditioned analysis reveals partial asymmetry. Under DistilBERT, *nurse|F* and *nurse|M* both retain 5 shared tokens (ratio 0.33 each), showing no directional asymmetry. Under RoBERTa, *nurse|F* retains 7 shared tokens (ratio 0.54) versus *nurse|M* retaining 6 (ratio 0.43), a modest asymmetry suggesting the learned feature space for *nurse* is somewhat more anchored in female-typical contextual patterns. This attribution-level asymmetry aligns with the fairness results in Section 6.2, where *nurse* shows one of the largest EO gap reductions under masking—suggesting that professions with unstable proxy representations are precisely those where masking provides the most benefit.

Qualitative Illustrations

The following three examples from the DistilBERT unmasked condition provide concrete illustration of the aggregate patterns reported above.

Example 1: Gender asymmetry in attributed tokens (*nurse*, unmasked). Both biographies are correctly predicted as *nurse*.

Male biography excerpt	Female biography excerpt
“He holds a Masters of Nursing Degree and has served in the Australian Defense force in various roles for over 32 years. He has presented in this field for over 20 years. . .”	“Her scope of practice encompasses health promotion, disease prevention, diagnosis and management of common and complex health care problems. . . throughout the aging process. . .”
Top tokens: <i>nursing</i> (3.18), <i>masters</i> (0.61), <i>presented</i> (0.59), <i>australian</i> (0.44)	Top tokens: <i>scope</i> (0.74), <i>health</i> (0.66), <i>aging</i> (0.61), <i>practice</i> (0.55), <i>her</i> (0.49)

The male prediction anchors on explicit credentials (*nursing*, *masters*); the female prediction distributes across social-care context words (*scope*, *aging*)—consistent with the *nurse|M* faithfulness asymmetry (0.756 vs. 0.354).

Example 2: Prediction flip after masking (backfire).

“[MASK] is a parent of a child with autism, and is dedicated to providing professional assistance to other parents. . . [MASK] can be

reached at (408) 843-7997, and is available for paid phone consultation. . . ”

True label: *psychologist*. Unmasked prediction: *psychologist* — top tokens: *parent* (1.15), *autism* (0.81), *child* (0.39), *professional* (0.33). Masked prediction: *attorney* × — top tokens: *408* (0.63), *autism* (0.34), *reached* (0.33), *parents* (0.28).

After masking *He*, the parent/child anchor loses salience; the phone number and “consultation” framing pattern-match *attorney*, producing an error rather than a fairness gain.

Example 3 (mini): Indirect proxy tokens driving a physician prediction.

“Ms. Croteau *practices medicine* in New Britain, CT and specializes in Allergy & Immunology. Ms. Croteau is affiliated with *Hospital* of Central Connecticut.”

Predicted: *physician* — top tokens: *medicine* (1.13), *practices* (0.96), *hospital* (0.52). No gender tokens appear in the top 5. Clinical role words alone drive the prediction, consistent with the interpretation that indirect contextual signals, not explicit gender markers, carry the primary predictive load.

Faithfulness Validation

While Integrated Gradients provides correlational insights into feature importance, the top- K erasure test provides intervention-based validation: erasing attributed tokens and measuring the resulting confidence drop assesses whether those tokens are functionally important for the prediction. Table 8 reports mean faithfulness scores by profession across both models and masking conditions. Across all 3,000 examples, 89.6% (DistilBERT) and 82.3% (RoBERTa) of predictions show positive faithfulness, consistent with the attributed tokens being genuinely load-bearing.

Faithfulness scores align with the proxy stability analysis: professions with lower shared ratios tend toward higher faithfulness (*dietitian*: 0.619, *nurse*: 0.395), while stable professions show lower faithfulness (*professor*: 0.160, *attorney*: 0.162). This supports the interpretation that attribution instability identifies professions whose predictions are concentrated in a small, functionally decisive feature set. Notably, the *model* profession under RoBERTa unmasked achieves the lowest faithfulness of all 20 professions (0.100), suggesting its predictions are

Profession	DistilBERT		RoBERTa	
	U	M	U	M
Dietitian	0.619	0.743	0.548	0.372
Nurse	0.395	0.518	0.203	0.197
Psychologist	0.356	0.378	0.232	0.252
Painter	0.473	0.418	0.339	0.342
Physician	0.233	0.191	0.235	0.200
Attorney	0.162	0.221	0.175	0.158
Professor	0.160	0.214	0.162	0.182
Software Engineer	0.171	0.204	0.193	0.199
Overall	0.330	0.364	0.243	0.249

Table 8: Mean faithfulness (probability drop after erasing top-5 attributed tokens). U=unmasked; M=masked. Higher = more functionally important tokens under erasure. The *Overall* row is the macro-average across all 20 professions (equal weight per profession, not per example).

distributed across many weakly-attributing tokens rather than a small decisive set.

Gender-conditioned faithfulness reveals a striking asymmetry for *nurse* under DistilBERT: *nurse*|M = 0.756 versus *nurse*|F = 0.354. Erasing the top-5 tokens for male nurse predictions nearly halves the model’s confidence, consistent with male nurse representations being concentrated in a critically small feature set. Under RoBERTa the asymmetry is also present: *nurse*|M = 0.534 versus *nurse*|F = 0.163. Full gender-conditioned faithfulness results are reported in Appendix D.

Masking increases overall faithfulness for both models (DistilBERT: +0.034; RoBERTa: +0.006), suggesting that after gender signals are removed, the remaining attributed tokens become more functionally decisive — consistent with the model concentrating predictions on a smaller, more critical feature set.

7 Conclusion

Fine-tuned models substantially outperform zero/few-shot baselines in both performance and fairness; scaling without supervision amplifies bias; masking reduces aggregate EO gaps in encoder models (by 0.017–0.019) but produces profession-level backfire effects in Pythia; and bias is encoded primarily through indirect contextual tokens whose importance is model-dependent and whose functional load is supported by erasure faithfulness testing. Notably, zero-shot EO gap grows from ~ 0.09 at 160M to ~ 0.21 at 1.4B parameters—suggesting that model capacity alone

is not a substitute for task-specific fine-tuning or explicit debiasing interventions. Gender-conditioned analysis further reveals that male nurse predictions are disproportionately dependent on a critically small feature set, with faithfulness nearly double that of female nurse predictions (0.756 vs. 0.354 under DistilBERT).

Limitations. The dataset provides binary gender labels inferred from pronoun usage by the original authors [De-Arteaga et al., 2019], which do not capture the full spectrum of gender identity and limit fairness evaluation to a male/female binary. Attribution is restricted to encoder models. Per-profession evaluation uses proportional stratification from the 3,000-sample canonical subset, resulting in small sample sizes for low-frequency professions (e.g. comedian $n = 27$, dietitian $n = 30$); per-profession fairness estimates for these occupations should be interpreted with caution.

Future work. Remaining directions include a formal cross-gender proxy bias score and extending attribution to generative models.

A Per-Occupation Proxy Word Tables

Top-5 attributed tokens (weighted score = $\text{count_topk} \times \text{mean attribution}$) per profession for DistilBERT and RoBERTa under unmasked (U) and masked (M) conditions. Aggregated over 3,000 test examples; stopwords and tokens shorter than 3 characters excluded. *Note: some RoBERTa entries (e.g. olog, ession, ition) reflect BPE subword fragments and should be interpreted as components of larger lexical units.*

Accountant

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
accounting	20.8	accounting	17.8
financial	9.5	financial	9.2
accounts	3.9	account	7.0
finance	3.1	business	6.4
tax	2.1	tax	5.6
<i>Masked</i>			
accounting	21.4	accounting	18.3
financial	10.6	financial	10.3
accounts	4.0	account	9.0
finance	4.0	business	8.2
tax	2.2	finance	5.1

Architect

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
architecture	29.4	architecture	43.5
architectural	15.7	design	21.3
design	13.0	architect	20.6
architect	11.1	architectural	13.0
architects	7.9	architects	9.0
<i>Masked</i>			
architecture	27.8	architecture	38.7
architectural	15.1	design	20.4
design	13.6	architects	13.3
architects	8.8	architectural	12.1
construction	4.1	architect	11.9

Attorney

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
law	53.3	law	183.4
attorney	37.1	attorney	54.9
legal	26.9	litigation	51.5
litigation	25.0	legal	46.3
clients	19.7	practice	42.9
<i>Masked</i>			
law	56.9	law	168.8
legal	28.6	legal	56.7
litigation	25.3	litigation	53.8
clients	19.7	practice	31.4
counsel	18.8	court	21.9

Comedian

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
comedy	16.4	comedy	16.7
stand	4.3	stand	6.5
comedians	2.9	comedians	5.8
performed	2.0	performed	2.1
–	–	show	1.5
<i>Masked</i>			
comedy	16.7	comedy	19.0
stand	3.8	comedians	6.1
comedians	3.1	stand	4.9
funny	2.4	show	1.4
performed	1.8	–	–

Composer

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
music	23.9	music	28.6
composition	6.4	composition	7.2
composer	6.3	scores	7.1
musical	5.0	musical	6.8
composing	4.2	compositions	6.7
<i>Masked</i>			
music	23.9	music	27.3
composition	7.2	composition	10.1
compositions	5.3	scores	6.6
musical	5.0	musical	6.0
composing	3.9	composing	5.6

Dentist

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
dental	105.5	ental	90.5
dentistry	48.8	dent	45.7
dentist	28.3	dental	33.4
bds	23.2	clinic	22.7
clinic	9.6	patients	21.1
<i>Masked</i>			
dental	125.2	ental	94.5
dentistry	49.9	dental	39.8
bds	18.7	dent	33.0
clinic	11.4	clinic	23.0
patients	9.2	patients	19.4

Dietitian

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
nutrition	37.4	nutrition	24.7
dietitian	35.1	diet	4.9
registered	6.4	ession	4.3
nutritional	3.9	nutritional	3.8
diet	3.1	health	2.2
<i>Masked</i>			
nutrition	57.3	nutrition	20.8
having	4.4	ession	7.1
experiences	3.8	diet	4.5
nutritional	3.5	nutritional	4.0
diet	3.4	specializes	2.7

Filmmaker

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
film	20.7	film	26.4
films	12.2	films	24.9
documentary	8.5	feature	7.8
filmmaking	7.8	filmmaker	7.2
directed	4.5	filmmaking	6.9
<i>Masked</i>			
film	21.4	film	27.4
films	11.7	films	20.9
filmmaking	7.2	filmmaking	6.6
documentary	5.7	documentary	5.8
directed	5.2	festivals	5.4

Journalist

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
journalism	34.5	journalism	32.0
journalist	17.9	news	19.6
news	12.9	journalist	15.7
editor	11.4	editor	13.7
journalists	8.4	reporting	8.9
<i>Masked</i>			
journalism	35.3	journalism	36.1
news	14.4	news	23.4
editor	12.5	editor	10.8
journalists	7.8	times	10.4
reporting	7.4	reporting	10.0

Model Note: freeones is a website name from the training corpus, not a profession descriptor — an example of dataset artefacts appearing in proxy tables.

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
model	16.4	model	28.1
freeones	15.2	born	20.8
modeling	14.0	modeling	13.9
born	7.1	currently	6.7
gallery	5.3	fashion	6.5
<i>Masked</i>			
freeones	22.7	modeling	17.7
modeling	15.2	gallery	8.5
born	8.0	modelling	8.4
gallery	5.8	currently	8.4
modelling	4.6	ranked	6.3

Nurse

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
nurse	127.3	nursing	91.5
nursing	54.6	nurse	76.2
practitioner	25.1	specialists	24.2
graduated	20.5	graduated	23.2
experiences	14.5	medical	21.8
<i>Masked</i>			
nursing	66.6	nursing	107.5
graduated	23.0	medical	20.3
nurse	8.9	doctors	17.4
nurses	7.6	ition	16.9
medicine	7.2	health	15.6

Painter

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
paintings	25.4	painting	41.5
painting	21.1	paintings	38.8
gallery	8.1	art	9.6
art	7.8	paint	8.6
paint	4.1	gallery	7.2
<i>Masked</i>			
paintings	23.8	painting	36.7
painting	22.1	paintings	35.4
art	8.7	art	11.9
gallery	7.1	paint	8.6
paint	4.1	gallery	6.2

Photographer

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
photography	146.1	photography	164.7
photographs	25.5	photographs	32.9
photographer	22.9	photographer	26.8
photographic	15.9	images	18.9
photo	14.3	photographic	17.8
<i>Masked</i>			
photography	153.9	photography	176.8
photographs	24.4	photographs	30.8
photo	17.8	images	17.1
photographic	17.0	photo	17.1
images	10.1	photographic	16.5

Physician

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
<i>medicine</i>	262.6	<i>medicine</i>	203.5
<i>practices</i>	77.5	<i>practices</i>	172.6
<i>medical</i>	64.1	<i>medical</i>	110.4
<i>physician</i>	36.1	<i>specializes</i>	45.8
<i>graduate</i>	16.7	<i>affiliated</i>	41.7
<i>Masked</i>			
<i>medicine</i>	238.6	<i>practices</i>	148.7
<i>practices</i>	101.1	<i>medicine</i>	121.6
<i>medical</i>	63.0	<i>medical</i>	106.7
<i>specializes</i>	29.1	<i>affiliated</i>	79.5
<i>graduate</i>	17.0	<i>specializes</i>	46.4

Poet

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
<i>poetry</i>	34.5	<i>poetry</i>	40.4
<i>poems</i>	8.1	<i>poems</i>	19.7
<i>writing</i>	4.4	<i>writing</i>	10.0
<i>english</i>	3.0	<i>poem</i>	8.3
<i>poem</i>	3.0	<i>published</i>	6.2
<i>Masked</i>			
<i>poetry</i>	35.9	<i>poetry</i>	47.3
<i>poems</i>	8.3	<i>poems</i>	18.8
<i>writing</i>	4.7	<i>writing</i>	13.5
<i>poem</i>	3.5	<i>etry</i>	6.3
<i>english</i>	3.2	<i>poem</i>	6.1

Professor

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
<i>research</i>	322.9	<i>research</i>	364.0
<i>phd</i>	92.0	<i>phd</i>	127.4
<i>interests</i>	58.8	<i>university</i>	126.5
<i>university</i>	58.6	<i>interests</i>	83.1
<i>focuses</i>	48.0	<i>teaching</i>	67.5
<i>Masked</i>			
<i>research</i>	220.1	<i>research</i>	302.4
<i>phd</i>	69.9	<i>university</i>	135.8
<i>university</i>	66.5	<i>phd</i>	100.4
<i>interests</i>	58.0	<i>interests</i>	68.0
<i>focuses</i>	50.2	<i>faculty</i>	63.5

Psychologist

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
<i>psychologist</i>	58.3	<i>psychology</i>	35.9
<i>psychology</i>	42.0	<i>olog</i>	30.0
<i>psychological</i>	9.1	<i>clinical</i>	13.2
<i>psychotherapy</i>	9.1	<i>psych</i>	12.3
<i>experiences</i>	8.0	<i>psychological</i>	10.5
<i>Masked</i>			
<i>psychology</i>	42.1	<i>psychology</i>	36.8
<i>psychological</i>	16.3	<i>psychological</i>	18.3
<i>psychotherapy</i>	9.5	<i>psych</i>	11.7
<i>depression</i>	9.2	<i>clinical</i>	11.1
<i>therapy</i>	6.4	<i>ical</i>	10.8

Software Engineer

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
<i>development</i>	5.2	<i>development</i>	9.3
<i>engineering</i>	4.6	<i>engineering</i>	9.2
<i>software</i>	4.2	<i>computer</i>	7.0
<i>ibm</i>	2.9	<i>software</i>	5.0
<i>developer</i>	2.7	<i>java</i>	3.7
<i>Masked</i>			
<i>engineering</i>	5.5	<i>development</i>	12.3
<i>development</i>	4.8	<i>engineering</i>	7.8
<i>software</i>	3.9	<i>software</i>	5.2
<i>ibm</i>	2.9	<i>computer</i>	4.5
<i>code</i>	2.4	<i>building</i>	3.2

Surgeon

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
<i>surgery</i>	43.7	<i>surgery</i>	83.7
<i>medical</i>	13.2	<i>medical</i>	27.8
<i>medicine</i>	11.7	<i>surgeon</i>	12.5
<i>practices</i>	8.5	<i>hospital</i>	9.0
<i>surgeon</i>	7.7	<i>surgeries</i>	8.9
<i>Masked</i>			
<i>surgery</i>	42.6	<i>surgery</i>	80.4
<i>practices</i>	10.5	<i>medical</i>	33.2
<i>medicine</i>	8.9	<i>practices</i>	12.9
<i>medical</i>	8.6	<i>hospital</i>	11.0
<i>practice</i>	6.8	<i>practice</i>	10.1

Teacher

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
<i>teacher</i>	19.5	<i>teaching</i>	26.9
<i>students</i>	18.2	<i>school</i>	25.3
<i>teaching</i>	15.8	<i>students</i>	25.2
<i>school</i>	12.2	<i>teacher</i>	23.0
<i>teachers</i>	12.1	<i>taught</i>	17.6
<i>Masked</i>			
<i>students</i>	19.4	<i>school</i>	32.9
<i>teachers</i>	15.0	<i>teaching</i>	26.6
<i>teaching</i>	14.5	<i>students</i>	26.2
<i>school</i>	11.3	<i>teachers</i>	15.8
<i>schools</i>	8.4	<i>teaches</i>	12.4

B Prompt Templates

B.1 Zero-Shot

```
""
Given the short biography below, predict the occupation of the person.
Answer with a single specific occupation word (e.g. professor,
physician, attorney, photographer, journalist, nurse, psychologist,
teacher, dentist, surgeon, architect, painter, model, poet, filmmaker,
software_engineer, accountant, composer, dietitian, comedian etc.).
```

```
Biography: {text}
Occupation:
""
```

B.2 Few-Shot

```
""
Given the short biography below, predict the occupation of the person.
Answer with a single specific occupation word (e.g. professor, physician,
attorney, photographer, journalist, nurse, psychologist, teacher, dentist,
surgeon, architect, painter, model, poet, filmmaker, software_engineer,
accountant, composer, dietitian, comedian etc.).
```

```
Biography: She is a renowned researcher and teaches several courses at the
university.
Occupation: professor
```

```
Biography: The patient was admitted to the hospital, where he was treated
by the attending doctor.
Occupation: physician
```

```
Biography: He covers local news and politics for the daily newspaper.
Occupation: journalist
```

```
Biography: He is a talented photographer who has won several awards for his
work.
Occupation: photographer
```

```
Biography: {text}
Occupation:
""
```

C Detailed Fairness Metrics by Occupation

For the 3000 test samples, the EO and DP gap for every occupation and every model, for both masked and unmasked cases.

Accountant

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.094	0.002	0.012	Pythia-410m FS	U	0.167	0.002	0.009
DistilBERT FT	M	0.031	0.000	0.009	Pythia-410m FS	M	0.167	0.002	0.009
RoBERTa FT	U	0.062	0.000	0.009	Pythia-410m 0s	U	0.073	0.001	0.004
RoBERTa FT	M	0.010	0.001	0.008	Pythia-410m 0s	M	0.073	0.000	0.003
Pythia-1.4B FT	U	0.229	0.002	0.013	Pythia-160m FT	U	0.000	0.000	0.000
Pythia-1.4B FT	M	0.229	0.002	0.013	Pythia-160m FT	M	0.000	0.000	0.000
Pythia-1.4B FS	U	0.333	0.004	0.015	Pythia-160m FS	U	0.083	0.000	0.001
Pythia-1.4B FS	M	0.333	0.005	0.016	Pythia-160m FS	M	0.083	0.000	0.001
Pythia-1.4B 0s	U	0.260	0.001	0.001	Pythia-160m 0s	U	0.000	0.000	0.000
Pythia-1.4B 0s	M	0.260	0.001	0.001	Pythia-160m 0s	M	0.000	0.000	0.000
Pythia-410m FT	U	0.062	0.003	0.001					
Pythia-410m FT	M	0.042	0.001	0.003					

Architect

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.242	0.006	0.017	Pythia-410m FS	U	0.118	0.010	0.021
DistilBERT FT	M	0.134	0.008	0.019	Pythia-410m FS	M	0.170	0.010	0.022
RoBERTa FT	U	0.173	0.008	0.020	Pythia-410m 0s	U	0.082	0.059	0.076
RoBERTa FT	M	0.156	0.007	0.020	Pythia-410m 0s	M	0.099	0.060	0.077
Pythia-1.4B FT	U	0.209	0.005	0.017	Pythia-160m FT	U	0.000	0.000	0.000
Pythia-1.4B FT	M	0.175	0.003	0.016	Pythia-160m FT	M	0.000	0.002	0.002
Pythia-1.4B FS	U	0.012	0.011	0.021	Pythia-160m FS	U	0.131	0.002	0.008
Pythia-1.4B FS	M	0.012	0.010	0.020	Pythia-160m FS	M	0.114	0.001	0.006
Pythia-1.4B 0s	U	0.135	0.002	0.005	Pythia-160m 0s	U	0.014	0.001	0.002
Pythia-1.4B 0s	M	0.135	0.002	0.005	Pythia-160m 0s	M	0.050	0.001	0.003
Pythia-410m FT	U	0.182	0.015	0.015					
Pythia-410m FT	M	0.028	0.002	0.005					

Attorney

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.046	0.001	0.019	Pythia-410m FS	U	0.153	0.004	0.018
DistilBERT FT	M	0.002	0.000	0.016	Pythia-410m FS	M	0.147	0.004	0.017
RoBERTa FT	U	0.026	0.000	0.018	Pythia-410m 0s	U	0.202	0.008	0.037
RoBERTa FT	M	0.030	0.003	0.016	Pythia-410m 0s	M	0.185	0.009	0.036
Pythia-1.4B FT	U	0.014	0.004	0.014	Pythia-160m FT	U	0.022	0.000	0.003
Pythia-1.4B FT	M	0.020	0.001	0.017	Pythia-160m FT	M	0.002	0.002	0.002
Pythia-1.4B FS	U	0.018	0.008	0.007	Pythia-160m FS	U	0.053	0.001	0.004
Pythia-1.4B FS	M	0.018	0.008	0.007	Pythia-160m FS	M	0.046	0.001	0.003
Pythia-1.4B 0s	U	0.065	0.005	0.001	Pythia-160m 0s	U	0.005	0.000	0.002
Pythia-1.4B 0s	M	0.065	0.005	0.001	Pythia-160m 0s	M	0.012	0.001	0.003
Pythia-410m FT	U	0.009	0.005	0.015					
Pythia-410m FT	M	0.042	0.004	0.010					

Comedian

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.059	0.004	0.010	Pythia-410m FS	U	0.271	0.002	0.006
DistilBERT FT	M	0.118	0.002	0.008	Pythia-410m FS	M	0.271	0.001	0.005
RoBERTa FT	U	0.141	0.003	0.010	Pythia-410m 0s	U	0.271	0.001	0.005
RoBERTa FT	M	0.141	0.002	0.009	Pythia-410m 0s	M	0.271	0.001	0.005
Pythia-1.4B FT	U	0.141	0.001	0.008	Pythia-160m FT	U	0.000	0.000	0.000
Pythia-1.4B FT	M	0.141	0.000	0.007	Pythia-160m FT	M	0.000	0.000	0.000
Pythia-1.4B FS	U	0.365	0.001	0.007	Pythia-160m FS	U	0.200	0.001	0.000
Pythia-1.4B FS	M	0.365	0.001	0.007	Pythia-160m FS	M	0.200	0.000	0.001
Pythia-1.4B 0s	U	0.047	0.000	0.002	Pythia-160m 0s	U	0.235	0.000	0.002
Pythia-1.4B 0s	M	0.047	0.000	0.002	Pythia-160m 0s	M	0.235	0.000	0.002
Pythia-410m FT	U	0.000	0.002	0.002					
Pythia-410m FT	M	0.000	0.002	0.002					

Composer

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.025	0.001	0.015	Pythia-410m FS	U	0.152	0.003	0.011
DistilBERT FT	M	0.025	0.001	0.015	Pythia-410m FS	M	0.152	0.005	0.013
RoBERTa FT	U	0.003	0.000	0.013	Pythia-410m 0s	U	0.003	0.002	0.016
RoBERTa FT	M	0.054	0.001	0.014	Pythia-410m 0s	M	0.003	0.003	0.016
Pythia-1.4B FT	U	0.025	0.001	0.015	Pythia-160m FT	U	0.000	0.000	0.000
Pythia-1.4B FT	M	0.054	0.000	0.015	Pythia-160m FT	M	0.000	0.000	0.000
Pythia-1.4B FS	U	0.162	0.006	0.019	Pythia-160m FS	U	0.051	0.001	0.003
Pythia-1.4B FS	M	0.162	0.006	0.019	Pythia-160m FS	M	0.051	0.001	0.003
Pythia-1.4B 0s	U	0.067	0.003	0.011	Pythia-160m 0s	U	0.051	0.001	0.002
Pythia-1.4B 0s	M	0.067	0.003	0.011	Pythia-160m 0s	M	0.022	0.001	0.002
Pythia-410m FT	U	0.165	0.004	0.004					
Pythia-410m FT	M	0.025	0.001	0.002					

Dentist

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.057	0.002	0.030	Pythia-410m FS	U	0.016	0.000	0.026
DistilBERT FT	M	0.080	0.002	0.029	Pythia-410m FS	M	0.016	0.001	0.027
RoBERTa FT	U	0.068	0.002	0.029	Pythia-410m 0s	U	0.006	0.002	0.026
RoBERTa FT	M	0.057	0.002	0.030	Pythia-410m 0s	M	0.006	0.002	0.025
Pythia-1.4B FT	U	0.068	0.002	0.029	Pythia-160m FT	U	0.000	0.000	0.000
Pythia-1.4B FT	M	0.068	0.002	0.029	Pythia-160m FT	M	0.000	0.000	0.000
Pythia-1.4B FS	U	0.045	0.004	0.031	Pythia-160m FS	U	0.201	0.001	0.021
Pythia-1.4B FS	M	0.045	0.004	0.030	Pythia-160m FS	M	0.198	0.001	0.024
Pythia-1.4B 0s	U	0.019	0.001	0.028	Pythia-160m 0s	U	0.016	0.000	0.016
Pythia-1.4B 0s	M	0.019	0.000	0.027	Pythia-160m 0s	M	0.027	0.001	0.016
Pythia-410m FT	U	0.082	0.002	0.019					
Pythia-410m FT	M	0.053	0.001	0.020					

Dietitian

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.103	0.000	0.018	Pythia-410m FS	U	0.552	0.002	0.010
DistilBERT FT	M	0.034	0.000	0.019	Pythia-410m FS	M	0.552	0.002	0.010
RoBERTa FT	U	0.069	0.001	0.019	Pythia-410m 0s	U	0.586	0.001	0.013
RoBERTa FT	M	0.103	0.000	0.018	Pythia-410m 0s	M	0.552	0.001	0.012
Pythia-1.4B FT	U	0.069	0.001	0.019	Pythia-160m FT	U	0.190	0.002	0.003
Pythia-1.4B FT	M	0.069	0.000	0.018	Pythia-160m FT	M	0.017	0.011	0.013
Pythia-1.4B FS	U	0.103	0.010	0.028	Pythia-160m FS	U	0.000	0.003	0.003
Pythia-1.4B FS	M	0.103	0.010	0.028	Pythia-160m FS	M	0.000	0.004	0.004
Pythia-1.4B 0s	U	0.862	0.000	0.018	Pythia-160m 0s	U	0.000	0.001	0.001
Pythia-1.4B 0s	M	0.862	0.000	0.018	Pythia-160m 0s	M	0.000	0.001	0.001
Pythia-410m FT	U	0.034	0.006	0.007					
Pythia-410m FT	M	0.034	0.003	0.004					

Filmmaker

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.079	0.003	0.007	Pythia-410m FS	U	0.071	0.003	0.002
DistilBERT FT	M	0.021	0.003	0.009	Pythia-410m FS	M	0.071	0.003	0.002
RoBERTa FT	U	0.064	0.003	0.008	Pythia-410m 0s	U	0.236	0.001	0.006
RoBERTa FT	M	0.043	0.003	0.009	Pythia-410m 0s	M	0.157	0.001	0.005
Pythia-1.4B FT	U	0.036	0.001	0.008	Pythia-160m FT	U	0.000	0.001	0.001
Pythia-1.4B FT	M	0.014	0.002	0.009	Pythia-160m FT	M	0.000	0.001	0.001
Pythia-1.4B FS	U	0.093	0.001	0.004	Pythia-160m FS	U	0.114	0.000	0.001
Pythia-1.4B FS	M	0.143	0.002	0.004	Pythia-160m FS	M	0.086	0.000	0.001
Pythia-1.4B 0s	U	0.071	0.000	0.001	Pythia-160m 0s	U	0.029	0.000	0.001
Pythia-1.4B 0s	M	0.071	0.000	0.001	Pythia-160m 0s	M	0.000	0.000	0.000
Pythia-410m FT	U	0.029	0.003	0.003					
Pythia-410m FT	M	0.029	0.002	0.002					

Journalist

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.076	0.003	0.012	Pythia-410m FS	U	0.106	0.037	0.047
DistilBERT FT	M	0.063	0.002	0.013	Pythia-410m FS	M	0.092	0.025	0.035
RoBERTa FT	U	0.111	0.001	0.019	Pythia-410m 0s	U	0.025	0.018	0.025
RoBERTa FT	M	0.048	0.004	0.010	Pythia-410m 0s	M	0.047	0.021	0.027
Pythia-1.4B FT	U	0.055	0.002	0.013	Pythia-160m FT	U	0.000	0.000	0.000
Pythia-1.4B FT	M	0.033	0.007	0.006	Pythia-160m FT	M	0.000	0.000	0.000
Pythia-1.4B FS	U	0.103	0.036	0.030	Pythia-160m FS	U	0.021	0.075	0.077
Pythia-1.4B FS	M	0.090	0.034	0.028	Pythia-160m FS	M	0.007	0.085	0.086
Pythia-1.4B 0s	U	0.350	0.004	0.028	Pythia-160m 0s	U	0.137	0.002	0.011
Pythia-1.4B 0s	M	0.350	0.004	0.028	Pythia-160m 0s	M	0.074	0.002	0.008
Pythia-410m FT	U	0.009	0.002	0.003					
Pythia-410m FT	M	0.043	0.006	0.002					

Model

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.173	0.002	0.028	Pythia-410m FS	U	0.127	0.014	0.001
DistilBERT FT	M	0.304	0.000	0.028	Pythia-410m FS	M	0.127	0.014	0.001
RoBERTa FT	U	0.153	0.002	0.023	Pythia-410m 0s	U	0.227	0.001	0.011
RoBERTa FT	M	0.062	0.001	0.027	Pythia-410m 0s	M	0.227	0.001	0.011
Pythia-1.4B FT	U	0.173	0.001	0.027	Pythia-160m FT	U	0.000	0.001	0.001
Pythia-1.4B FT	M	0.082	0.001	0.027	Pythia-160m FT	M	0.000	0.001	0.001
Pythia-1.4B FS	U	0.233	0.004	0.031	Pythia-160m FS	U	0.069	0.016	0.012
Pythia-1.4B FS	M	0.253	0.005	0.032	Pythia-160m FS	M	0.049	0.016	0.013
Pythia-1.4B 0s	U	0.067	0.004	0.009	Pythia-160m 0s	U	0.162	0.002	0.002
Pythia-1.4B 0s	M	0.067	0.004	0.009	Pythia-160m 0s	M	0.031	0.002	0.000
Pythia-410m FT	U	0.209	0.003	0.008					
Pythia-410m FT	M	0.380	0.003	0.010					

Nurse

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.060	0.012	0.086	Pythia-410m FS	U	0.179	0.029	0.085
DistilBERT FT	M	0.074	0.006	0.078	Pythia-410m FS	M	0.235	0.028	0.085
RoBERTa FT	U	0.067	0.009	0.084	Pythia-410m 0s	U	0.056	0.020	0.084
RoBERTa FT	M	0.019	0.011	0.082	Pythia-410m 0s	M	0.034	0.021	0.083
Pythia-1.4B FT	U	0.100	0.012	0.086	Pythia-160m FT	U	0.056	0.001	0.000
Pythia-1.4B FT	M	0.034	0.009	0.079	Pythia-160m FT	M	0.033	0.001	0.000
Pythia-1.4B FS	U	0.024	0.062	0.123	Pythia-160m FS	U	0.137	0.004	0.017
Pythia-1.4B FS	M	0.024	0.062	0.123	Pythia-160m FS	M	0.107	0.006	0.015
Pythia-1.4B 0s	U	0.660	0.014	0.085	Pythia-160m 0s	U	0.260	0.006	0.030
Pythia-1.4B 0s	M	0.660	0.015	0.086	Pythia-160m 0s	M	0.260	0.007	0.031
Pythia-410m FT	U	0.113	0.005	0.040					
Pythia-410m FT	M	0.232	0.004	0.037					

Painter

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.178	0.001	0.005	Pythia-410m FS	U	0.044	0.003	0.004
DistilBERT FT	M	0.046	0.002	0.000	Pythia-410m FS	M	0.025	0.003	0.002
RoBERTa FT	U	0.057	0.003	0.002	Pythia-410m 0s	U	0.030	0.000	0.001
RoBERTa FT	M	0.046	0.001	0.000	Pythia-410m 0s	M	0.004	0.001	0.000
Pythia-1.4B FT	U	0.043	0.000	0.002	Pythia-160m FT	U	0.000	0.001	0.001
Pythia-1.4B FT	M	0.078	0.001	0.001	Pythia-160m FT	M	0.000	0.000	0.000
Pythia-1.4B FS	U	0.001	0.000	0.000	Pythia-160m FS	U	0.018	0.001	0.001
Pythia-1.4B FS	M	0.001	0.000	0.000	Pythia-160m FS	M	0.053	0.002	0.001
Pythia-1.4B 0s	U	0.105	0.003	0.000	Pythia-160m 0s	U	0.065	0.017	0.015
Pythia-1.4B 0s	M	0.139	0.003	0.001	Pythia-160m 0s	M	0.096	0.020	0.017
Pythia-410m FT	U	0.059	0.003	0.004					
Pythia-410m FT	M	0.018	0.004	0.004					

Photographer

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.013	0.001	0.023	Pythia-410m FS	U	0.079	0.044	0.061
DistilBERT FT	M	0.001	0.002	0.022	Pythia-410m FS	M	0.065	0.038	0.055
RoBERTa FT	U	0.054	0.002	0.029	Pythia-410m 0s	U	0.040	0.004	0.026
RoBERTa FT	M	0.030	0.002	0.023	Pythia-410m 0s	M	0.031	0.003	0.024
Pythia-1.4B FT	U	0.065	0.000	0.026	Pythia-160m FT	U	0.014	0.002	0.003
Pythia-1.4B FT	M	0.028	0.000	0.025	Pythia-160m FT	M	0.000	0.000	0.000
Pythia-1.4B FS	U	0.019	0.014	0.033	Pythia-160m FS	U	0.084	0.003	0.013
Pythia-1.4B FS	M	0.033	0.015	0.035	Pythia-160m FS	M	0.123	0.001	0.018
Pythia-1.4B 0s	U	0.057	0.008	0.008	Pythia-160m 0s	U	0.062	0.000	0.003
Pythia-1.4B 0s	M	0.057	0.008	0.008	Pythia-160m 0s	M	0.035	0.001	0.004
Pythia-410m FT	U	0.046	0.005	0.024					
Pythia-410m FT	M	0.086	0.022	0.041					

Physician

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.087	0.006	0.027	Pythia-410m FS	U	0.055	0.033	0.020
DistilBERT FT	M	0.113	0.002	0.032	Pythia-410m FS	M	0.034	0.030	0.016
RoBERTa FT	U	0.087	0.001	0.031	Pythia-410m 0s	U	0.147	0.043	0.036
RoBERTa FT	M	0.076	0.003	0.028	Pythia-410m 0s	M	0.142	0.052	0.043
Pythia-1.4B FT	U	0.071	0.003	0.027	Pythia-160m FT	U	0.000	0.001	0.001
Pythia-1.4B FT	M	0.083	0.002	0.029	Pythia-160m FT	M	0.006	0.003	0.002
Pythia-1.4B FS	U	0.065	0.014	0.011	Pythia-160m FS	U	0.181	0.003	0.010
Pythia-1.4B FS	M	0.059	0.015	0.010	Pythia-160m FS	M	0.153	0.006	0.003
Pythia-1.4B 0s	U	0.164	0.016	0.021	Pythia-160m 0s	U	0.195	0.016	0.003
Pythia-1.4B 0s	M	0.151	0.016	0.020	Pythia-160m 0s	M	0.201	0.016	0.003
Pythia-410m FT	U	0.104	0.013	0.018					
Pythia-410m FT	M	0.129	0.015	0.018					

Poet

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.220	0.000	0.004	Pythia-410m FS	U	0.107	0.002	0.004
DistilBERT FT	M	0.153	0.002	0.004	Pythia-410m FS	M	0.107	0.004	0.006
RoBERTa FT	U	0.153	0.001	0.003	Pythia-410m 0s	U	0.147	0.000	0.003
RoBERTa FT	M	0.120	0.002	0.004	Pythia-410m 0s	M	0.113	0.001	0.001
Pythia-1.4B FT	U	0.120	0.001	0.003	Pythia-160m FT	U	0.000	0.000	0.000
Pythia-1.4B FT	M	0.187	0.002	0.005	Pythia-160m FT	M	0.000	0.000	0.000
Pythia-1.4B FS	U	0.153	0.001	0.003	Pythia-160m FS	U	0.060	0.003	0.004
Pythia-1.4B FS	M	0.153	0.000	0.003	Pythia-160m FS	M	0.060	0.004	0.005
Pythia-1.4B 0s	U	0.133	0.002	0.004	Pythia-160m 0s	U	0.020	0.008	0.009
Pythia-1.4B 0s	M	0.133	0.003	0.005	Pythia-160m 0s	M	0.093	0.008	0.010
Pythia-410m FT	U	0.060	0.001	0.002					
Pythia-410m FT	M	0.127	0.004	0.002					

Professor

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.024	0.012	0.022	Pythia-410m FS	U	0.040	0.006	0.018
DistilBERT FT	M	0.022	0.024	0.031	Pythia-410m FS	M	0.042	0.005	0.018
RoBERTa FT	U	0.018	0.024	0.032	Pythia-410m 0s	U	0.299	0.029	0.118
RoBERTa FT	M	0.020	0.019	0.028	Pythia-410m 0s	M	0.301	0.028	0.118
Pythia-1.4B FT	U	0.003	0.012	0.030	Pythia-160m FT	U	0.046	0.011	0.005
Pythia-1.4B FT	M	0.003	0.002	0.023	Pythia-160m FT	M	0.006	0.017	0.014
Pythia-1.4B FS	U	0.120	0.010	0.050	Pythia-160m FS	U	0.097	0.019	0.047
Pythia-1.4B FS	M	0.119	0.010	0.050	Pythia-160m FS	M	0.120	0.024	0.057
Pythia-1.4B 0s	U	0.078	0.243	0.206	Pythia-160m 0s	U	0.144	0.231	0.211
Pythia-1.4B 0s	M	0.083	0.242	0.207	Pythia-160m 0s	M	0.165	0.215	0.206
Pythia-410m FT	U	0.020	0.026	0.007					
Pythia-410m FT	M	0.020	0.026	0.008					

Psychologist

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.006	0.002	0.023	Pythia-410m FS	U	0.060	0.001	0.010
DistilBERT FT	M	0.067	0.003	0.021	Pythia-410m FS	M	0.067	0.001	0.010
RoBERTa FT	U	0.034	0.000	0.024	Pythia-410m 0s	U	0.024	0.004	0.014
RoBERTa FT	M	0.002	0.001	0.024	Pythia-410m 0s	M	0.013	0.004	0.015
Pythia-1.4B FT	U	0.060	0.002	0.025	Pythia-160m FT	U	0.011	0.002	0.002
Pythia-1.4B FT	M	0.031	0.002	0.024	Pythia-160m FT	M	0.000	0.001	0.001
Pythia-1.4B FS	U	0.010	0.023	0.038	Pythia-160m FS	U	0.026	0.005	0.005
Pythia-1.4B FS	M	0.010	0.023	0.038	Pythia-160m FS	M	0.008	0.003	0.004
Pythia-1.4B 0s	U	0.305	0.006	0.030	Pythia-160m 0s	U	0.003	0.000	0.002
Pythia-1.4B 0s	M	0.294	0.006	0.029	Pythia-160m 0s	M	0.028	0.000	0.002
Pythia-410m FT	U	0.175	0.011	0.025					
Pythia-410m FT	M	0.120	0.006	0.017					

Software Engineer

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.070	0.010	0.028	Pythia-410m FS	U	0.021	0.003	0.003
DistilBERT FT	M	0.070	0.007	0.025	Pythia-410m FS	M	0.043	0.003	0.004
RoBERTa FT	U	0.112	0.009	0.026	Pythia-410m 0s	U	0.155	0.006	0.013
RoBERTa FT	M	0.134	0.008	0.024	Pythia-410m 0s	M	0.112	0.006	0.012
Pythia-1.4B FT	U	0.027	0.009	0.029	Pythia-160m FT	U	0.000	0.001	0.001
Pythia-1.4B FT	M	0.036	0.007	0.029	Pythia-160m FT	M	0.000	0.001	0.001
Pythia-1.4B FS	U	0.322	0.040	0.061	Pythia-160m FS	U	0.021	0.002	0.003
Pythia-1.4B FS	M	0.322	0.040	0.061	Pythia-160m FS	M	0.021	0.004	0.005
Pythia-1.4B 0s	U	0.103	0.014	0.027	Pythia-160m 0s	U	0.000	0.001	0.001
Pythia-1.4B 0s	M	0.103	0.014	0.027	Pythia-160m 0s	M	0.000	0.001	0.001
Pythia-410m FT	U	0.106	0.001	0.004					
Pythia-410m FT	M	0.043	0.000	0.001					

Surgeon

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.070	0.010	0.041	Pythia-410m FS	U	0.045	0.002	0.010
DistilBERT FT	M	0.020	0.007	0.037	Pythia-410m FS	M	0.007	0.002	0.010
RoBERTa FT	U	0.043	0.011	0.042	Pythia-410m 0s	U	0.018	0.001	0.010
RoBERTa FT	M	0.018	0.007	0.037	Pythia-410m 0s	M	0.018	0.002	0.011
Pythia-1.4B FT	U	0.116	0.012	0.045	Pythia-160m FT	U	0.000	0.000	0.000
Pythia-1.4B FT	M	0.055	0.009	0.041	Pythia-160m FT	M	0.000	0.000	0.000
Pythia-1.4B FS	U	0.002	0.013	0.030	Pythia-160m FS	U	0.016	0.001	0.002
Pythia-1.4B FS	M	0.002	0.013	0.030	Pythia-160m FS	M	0.004	0.001	0.002
Pythia-1.4B 0s	U	0.116	0.001	0.000	Pythia-160m 0s	U	0.000	0.000	0.000
Pythia-1.4B 0s	M	0.104	0.001	0.001	Pythia-160m 0s	M	0.000	0.000	0.000
Pythia-410m FT	U	0.198	0.011	0.030					
Pythia-410m FT	M	0.252	0.005	0.024					

Teacher

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.073	0.000	0.021	Pythia-410m FS	U	0.157	0.033	0.045
DistilBERT FT	M	0.040	0.006	0.025	Pythia-410m FS	M	0.170	0.035	0.047
RoBERTa FT	U	0.132	0.010	0.033	Pythia-410m 0s	U	0.227	0.202	0.220
RoBERTa FT	M	0.086	0.004	0.025	Pythia-410m 0s	M	0.227	0.208	0.225
Pythia-1.4B FT	U	0.099	0.010	0.032	Pythia-160m FT	U	0.000	0.000	0.000
Pythia-1.4B FT	M	0.178	0.009	0.034	Pythia-160m FT	M	0.000	0.000	0.000
Pythia-1.4B FS	U	0.228	0.054	0.074	Pythia-160m FS	U	0.037	0.004	0.008
Pythia-1.4B FS	M	0.228	0.054	0.074	Pythia-160m FS	M	0.056	0.007	0.012
Pythia-1.4B 0s	U	0.534	0.060	0.091	Pythia-160m 0s	U	0.447	0.205	0.225
Pythia-1.4B 0s	M	0.534	0.061	0.092	Pythia-160m 0s	M	0.513	0.205	0.228
Pythia-410m FT	U	0.043	0.001	0.005					
Pythia-410m FT	M	0.053	0.010	0.017					

D Gender-Conditioned Faithfulness by Profession

Mean faithfulness scores split by gender for key professions, shown for all four conditions (two models $\times$ masked/unmasked). Only professions with at least 10 examples per gender group are shown; n = number of examples. Dietitian is excluded from masked rows (male $n=3$ in both masked conditions).

Profession	Female		Male	
	Faith.	n	Faith.	n
<i>DistilBERT — unmasked</i>				
Nurse	0.354	132	0.756	15
Dietitian	0.656	27	0.288	3
Psychologist	0.370	79	0.336	55
Physician	0.198	171	0.270	157
Attorney	0.161	110	0.162	160
Professor	0.158	413	0.162	519
<i>DistilBERT — masked</i>				
Nurse	0.494	126	0.660	22
Psychologist	0.391	69	0.359	47
Physician	0.183	173	0.201	150
Attorney	0.223	115	0.220	161
Professor	0.197	422	0.227	544
<i>RoBERTa — unmasked</i>				
Nurse	0.163	130	0.534	16
Dietitian	0.549	28	0.533	2
Psychologist	0.234	77	0.227	52
Physician	0.280	174	0.184	154
Attorney	0.192	108	0.164	156
Professor	0.155	411	0.168	533
<i>RoBERTa — masked</i>				
Nurse	0.172	129	0.374	18
Psychologist	0.274	80	0.221	55
Physician	0.218	170	0.181	154
Attorney	0.190	109	0.134	153
Professor	0.172	410	0.191	525

The nurse|M asymmetry is the largest gender-conditioned faithfulness gap observed in both unmasked conditions: DistilBERT 0.756 vs. 0.354, RoBERTa 0.534 vs. 0.163. Under masking, the gap narrows substantially (DistilBERT: 0.660 vs. 0.494; RoBERTa: 0.374 vs. 0.172) but does not disappear, consistent with implicit proxy signals persisting after explicit gender-token removal.

References

- Giuseppe Attanasio, Debora Nozza, Dirk Hovy, and Elena Baralis. Entropy-based attention regularization frees unintended bias mitigation from lists. In *Findings of the Association for Computational Linguistics: ACL 2022*, pages 1105–1119, 2022.
- Stella Biderman, Hailey Schoelkopf, Quentin Gregory Anthony, Herbie Bradley, Kyle O’Brien, Eric Hallahan, Mohammad Aflah Khan, Shivanshu Purohit, USVSN Sai Prashanth, Edward Raff, et al. Pythia: A suite for analyzing large language models across training and scaling. In *Proceedings of the 40th International Conference on Machine Learning*, pages 2397–2430, 2023.
- Rishi Bommasani, Drew A Hudson, Ehsan Adeli, et al. On the opportunities and risks of foundation models. *arXiv preprint arXiv:2108.07258*, 2021.
- Maria De-Arteaga, Alexey Romanov, Hanna Wallach, Jennifer Chayes, Christian Borgs, Alexandra Chouldechova, Sahin Geyik, Krishnaram Kenthapadi, and Adam Tauman Kalai. Bias in bios: A case study of semantic representation bias in a high-stakes setting. In *Proceedings of the ACM Conference on Fairness, Accountability, and Transparency*, pages 120–128, 2019.

- Cynthia Dwork, Moritz Hardt, Toniann Pitassi, Omer Reingold, and Richard Zemel. Fairness through awareness. *Proceedings of the 3rd Innovations in Theoretical Computer Science Conference*, pages 214–226, 2012.
- Moritz Hardt, Eric Price, and Nati Srebro. Equality of opportunity in supervised learning. In *Advances in Neural Information Processing Systems*, volume 29, 2016.
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. Lora: Low-rank adaptation of large language models. In *Proceedings of the International Conference on Learning Representations*, 2022.
- Yinhan Liu, Mili Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. Roberta: A robustly optimized bert pretraining approach. *arXiv preprint arXiv:1907.11692*, 2019.
- Kaiji Lu, Piotr Mardziel, Fangjing Wu, Preethi Amancharla, and Anupam Datta. Gender bias in neural natural language processing. In *Logic, Language, Information, and Computation*, pages 189–202, 2020.
- Alexey Romanov, Maria De-Arteaga, Hanna Wallach, Jennifer Chayes, Christian Borgs, Alexandra Chouldechova, Sahin Geyik, Krishnaram Kenthapadi, Anna Rumshisky, and Adam Tauman Kalai. What’s in a name? reducing bias in bios without access to protected attributes. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics*, pages 4195–4207, 2019.
- Victor Sanh, Lysandre Debut, Julien Chaumond, and Thomas Wolf. Distilbert, a distilled version of bert: smaller, faster, cheaper and lighter. *arXiv preprint arXiv:1910.01108*, 2019.
- Mukund Sundararajan, Ankur Taly, and Qiqi Yan. Axiomatic attribution for deep networks. In *Proceedings of the 34th International Conference on Machine Learning*, pages 3319–3328, 2017.