

Fairness and Proxy Bias in Occupation Prediction: Scaling, Masking, and Mechanistic Attribution across Model Families

Usukhbayar Purevdorj

Yitong Kong

Kirill Kamakaev

Giulio Zhu

Kian Makela

Ambrose Yang

University College London — COMP0087

Abstract

We investigate gender bias and proxy feature encoding in occupation prediction models trained on the Bias-in-Bios dataset. Our study spans two dimensions: *model capacity*, comparing Pythia generative models (160M–1.4B parameters) across zero-shot, few-shot, and fine-tuned regimes, and *model architecture*, contrasting encoder-based classifiers (DistilBERT, RoBERTa) against generative models of comparable performance. We evaluate both task performance (Macro-F1) and gender fairness (Equalized Odds gap, Demographic Parity) under masked and unmasked conditions, and diagnose bias mechanisms through Integrated Gradients attribution validated by top- K token erasure. Our main findings: (1) fine-tuned models substantially outperform zero/few-shot baselines in both performance and fairness; (2) larger models are not inherently fairer—zero-shot bias *increases* with scale; (3) gender-token masking reduces bias for some professions but produces backfire effects for others, suggesting that implicit proxy signals survive removal of explicit gender tokens; (4) explicit gender tokens are not the primary drivers of bias—they appear with weights 33–287 $\times$ smaller than dominant profession tokens, while indirect contextual signals carry the latent gender signal; and (5) proxy bias is model-dependent, the same profession can be stable under one model and unstable under another, though the precise cause (architecture, pre-training, or scale) cannot be isolated from this comparison alone.

1 Introduction

Occupation prediction from biographical text exposes persistent gender stereotypes: models trained on real-world career data reliably associate professions such as *nurse* or *dietitian* with female-typical language patterns, and *surgeon* or *software engineer* with male-typical ones. De-Arteaga et al.

[2019] demonstrated on the Bias-in-Bios dataset that scrubbing explicit gender signals does not eliminate these disparities—implicit proxy features survive.

This paper addresses three open questions: how does model *capacity* interact with fairness; how does model *architecture* shape which proxy features are learned; and what are those features precisely.

Contributions. *Primary contribution* (to our knowledge, limited prior work on this dataset provides systematic attribution analysis of this kind): (i) token-level Integrated Gradients attribution tables stratified by both profession and gender, with proxy stability and erasure faithfulness compared across encoder architectures. *Secondary contribution*: (ii) a cross-model comparison of how proxy token distributions shift under gender masking, quantifying both stability and functional importance. *Empirical applications of known methods to this setting*: (iii) a Pythia scaling study (160M–1.4B) across zero-shot, few-shot, and fine-tuned regimes; (iv) a cross-architecture comparison of encoder and generative model fairness on Bias-in-Bios; (v) a paired masked/unmasked evaluation of gender-token scrubbing on both task performance and fairness metrics.

Full findings are reported in §6.2 and §6.3. We do not claim these findings generalise beyond the Bias-in-Bios setting.

2 Related Work

Bias in occupation prediction. De-Arteaga et al. [2019] introduced Bias-in-Bios and showed that classifiers exhibit gender disparities even after pronoun scrubbing. Romanov et al. [2019] analysed limits of de-biasing through scrubbing. We extend these findings by identifying and quantifying specific proxy tokens through gradient-based attribution.

Fairness metrics. [Hardt et al. \[2016\]](#) formalised Equalized Odds; [Dwork et al. \[2012\]](#) introduced demographic parity. We evaluate both using a one-vs-rest multiclass extension macro-averaged over occupations.

Attribution methods. [Sundararajan et al. \[2017\]](#) introduced Integrated Gradients satisfying completeness and sensitivity axioms. Attribution methods have been applied to gender bias in coreference [[Attanasio et al., 2022](#)] but their systematic use in occupation classification is limited. Critically, no prior work on the Bias-in-Bios occupation prediction task provides token-level proxy word tables with quantified attribution weights, nor examines how those weights shift under gender masking.

Attribution method selection. Local surrogate methods such as LIME [[Ribeiro et al., 2016](#)] and additive feature attribution methods such as SHAP [[Lundberg and Lee, 2017](#)] offer alternative interpretability frameworks. LIME fits a locally linear model around each prediction and can be unstable across runs; KernelSHAP satisfies consistency and local accuracy axioms but requires a large number of model evaluations for each example, making it costly for token-level analysis of long sequences. Integrated Gradients [[Sundararajan et al., 2017](#)] provides a deterministic, path-integral attribution satisfying completeness (attributions sum to the prediction gap from the baseline) and sensitivity, with a fixed cost of $O(m)$ forward passes for m interpolation steps. For sequence classification tasks where token-level granularity and reproducibility matter, IG is the standard choice [[Attanasio et al., 2022](#)]; we adopt it for the same reasons.

Model scaling and fairness. [Bommasani et al. \[2021\]](#) discusses risks of foundation models including bias amplification. Our scaling study provides empirical evidence that zero-shot bias increases with model size in the Pythia family.

Gender masking. [Lu et al. \[2020\]](#) study counterfactual data augmentation for gender bias. Unlike prior work, we measure not only the effect of masking on fairness metrics but also on token-level attribution patterns — providing the first internal view of how occupation prediction models redistribute feature importance when gender signals are removed.

3 Dataset and Preprocessing

3.1 Bias-in-Bios

We use the Bias-in-Bios dataset [[De-Arteaga et al., 2019](#)] (LabHC/bias_in_bios), a widely-used benchmark in the NLP fairness and gender bias community, primarily employed to study occupational stereotypes and evaluate debiasing methods. Each example is a biography (`hard_text`) with an integer profession label and a binary gender attribute (inferred from pronoun usage by the original authors; $0 \rightarrow \text{M}$, $1 \rightarrow \text{F}$). We restrict to the **top 20 most frequent occupations** from the training split. The task is to predict occupation from biography text alone; gender is a post-hoc evaluation attribute only.

Split	Total	Female	Male
Train	248,141	46.16%	53.84%
Dev	38,198	46.16%	53.84%
Test	95,468	46.17%	53.83%
Total	381,807	46.16%	53.84%

Table 1: Dataset split sizes and gender distribution. The near-identical F% across splits (range: 0.01%, i.e. approximately 46.2% female in each) indicates the original dataset was stratified by gender when partitioned.

Table 2 shows the per-profession counts and gender distribution in the training split. Several occupations exhibit strong gender skew: *dietitian* (92.9% F), *nurse* (90.8% F), and *model* (82.7% F) are heavily female-dominated, while *surgeon* (14.8% F), *software engineer* (15.8% F), and *composer* (16.4% F) are heavily male-dominated.

3.2 Evaluation Subset

All models are evaluated on a canonical **3,000-sample test subset**, stratified proportionally by profession from the 95,468-example test split; a shared ID list ensures cross-model comparability. The subset was limited to 3,000 examples due to the computational cost of Integrated Gradients attribution, which requires 32 forward passes per example; running attribution on the full 95,468-example test split was not feasible within available compute. Per-profession sizes range from 932 (professor) to 27 (comedian); results for low-frequency professions should be interpreted with caution.

3.3 Gender Masking

Two conditions: **Unmasked** (original text) and **Masked** (gender tokens replaced with the model’s

Profession	N (train)	F%
professor	76,748	45.1
physician	26,648	49.4
attorney	21,169	38.3
photographer	15,773	35.7
journalist	12,960	49.5
nurse	12,316	90.8
psychologist	11,945	62.1
teacher	10,531	60.2
dentist	9,479	35.3
surgeon	8,829	14.8
architect	6,568	23.7
painter	5,025	45.7
model	4,867	82.7
filmmaker	4,545	32.9
poet	4,558	49.0
software engineer	4,492	15.8
composer	3,637	16.4
accountant	3,660	36.7
dietitian	2,567	92.9
comedian	1,824	21.1

Table 2: Per-profession training counts and F% (M% = 100 − F%). Distribution is consistent across all splits.

native mask token—[MASK] for DistilBERT, <mask> for RoBERTa—via case-insensitive word-boundary regex). Three token categories are masked: pronouns and possessives (*he, she, his, her, etc.*), titles and honorifics (*mr, mrs, ms, sir, lady, etc.*), and 54 gendered nouns (*father, mother, husband, wife, actor, actress, etc.*). Name masking was not applied: general name rules introduce uncontrolled false positives, and names carry biographical content beyond gender that would confound gender-signal removal with loss of biographical information.

4 Methods

4.1 Pythia: Zero-shot and Few-shot

We evaluate Pythia [Biderman et al., 2023] at 160M, 410M, and 1.4B parameters under zero-shot (biography + candidate list) and few-shot ($K = 3$ exemplar pairs) regimes, with and without gender masking.¹ As a causal language model, Pythia

¹The three exemplars cover professor (female pronoun), physician, and journalist (both male pronouns). To assess whether this 2:1 male skew introduced exemplar-level bias, we compared few-shot vs. zero-shot TPR gaps for these three professions across all three model sizes; no consistent directional effect was observed (full values in Appendix C), so exemplar gender is unlikely to be a primary confound in the few-shot results.

scores each candidate label c_k by its joint log-probability given prompt p :

$$\log P(c_k | p) = \sum_{t=1}^{|c_k|} \log P(c_k^{(t)} | p, c_k^{(1)}, \dots, c_k^{(t-1)}) \quad (1)$$

The predicted label is $\hat{y} = \arg \max_{c_k} \log P(c_k | p)$.

4.2 Pythia: Supervised Fine-tuning

Pythia is fine-tuned as a sequence classifier (linear head over the final hidden state) using LoRA [Hu et al., 2022] on attention projections ($r = 16$, $\alpha = 32$, dropout 0.1), lr 2×10^{-5} , batch 16, 3 epochs, max length 256; best checkpoint selected by dev Macro-F1. QLoRA (4-bit) is used for 1.4B.

4.3 Encoder Baselines: DistilBERT and RoBERTa

DistilBERT [Sanh et al., 2019] (66M) and RoBERTa-base [Liu et al., 2019] (125M) are fine-tuned as sequence classifiers with a linear head, lr 2×10^{-5} , batch 16, 3 epochs, max length 256, best checkpoint by Macro-F1. The same pipeline runs for both masked and unmasked conditions with a shared label mapping.

4.4 Evaluation Metrics

We evaluate models on both task performance and fairness. Let Y represent the true occupation label, $\hat{Y}$ the predicted label, $A \in \{M, F\}$ the binary gender attribute, and $\mathcal{C}$ the set of 20 occupation classes.

Task Performance: We report overall Accuracy and Macro-F1 score to account for the heavy class imbalance present in the dataset.

Fairness Metrics: We measure gender disparities using one-vs-rest multiclass extensions of Demographic Parity (DP) and Equalized Odds (EO), macro-averaged across all occupations.

Demographic Parity requires the prediction to be independent of the sensitive attribute. We compute the macro-averaged DP gap as the absolute difference in positive prediction rates between genders for each class:

$$\Delta DP = \frac{1}{|\mathcal{C}|} \sum_{c \in \mathcal{C}} \left| P(\hat{Y} = c | A = M) - P(\hat{Y} = c | A = F) \right|$$

Equalized Odds requires the prediction to be conditionally independent of the sensitive attribute

given the true label. To compute the EO gap, we first calculate the macro-averaged True Positive Rate (TPR) gap and False Positive Rate (FPR) gap:

$$\Delta\text{TPR} = \frac{1}{|\mathcal{C}|} \sum_{c \in \mathcal{C}} \left| P(\hat{Y} = c \mid Y = c, A = M) - P(\hat{Y} = c \mid Y = c, A = F) \right|$$

$$\Delta\text{FPR} = \frac{1}{|\mathcal{C}|} \sum_{c \in \mathcal{C}} \left| P(\hat{Y} = c \mid Y \neq c, A = M) - P(\hat{Y} = c \mid Y \neq c, A = F) \right|$$

The final Equalized Odds gap is defined as the maximum of these two macro-averaged disparities:

$$\text{EO Gap} = \max(\Delta\text{TPR}, \Delta\text{FPR})$$

4.5 Proxy Bias Attribution

Integrated Gradients [Sundararajan et al., 2017] is applied to the two fine-tuned encoder models, with attribution computed toward the *predicted* class index. Attributions are computed using **32 interpolation steps** and a **zero-embedding baseline**. Token-level scores are obtained by **summing** attribution values across the embedding dimension for each token position. Subword tokens are merged to word level before top- K selection: BERT-style ## continuation pieces and whitespace-prefixed tokens (RoBERTa’s byte-pair encoding marks word boundaries with a special prefix character) are handled separately, with scores summed across pieces. Only tokens with **positive** attribution scores and minimum length of 2 characters are retained. Maximum sequence length is **256 tokens**.

The pipeline operates in two stages with distinct K values. **Stage 1 — per-example extraction** ($K_1 = 5$): the 5 highest-scoring positive tokens are retained per example. **Stage 2 — profession-level aggregation**: tokens appearing across all examples for a profession are aggregated (minimum length 3, stopwords removed) to compute:

$$\text{wscore}(t, p) = \text{count_topk}(t, p) \times \overline{\text{attr}}(t, p) \quad (2)$$

where $\text{count_topk}(t, p)$ is the number of examples in which token t appears in the stage-1 top-5 for profession p , and $\overline{\text{attr}}(t, p)$ is the mean attribution score. Aggregation uses the model’s *predicted* label (not the true label): we attribute toward the predicted class to analyse the model’s decision mechanism rather than its correctness—the goal is to

understand what the model relies on, not whether it is right [Sundararajan et al., 2017].

Stability analysis. The masked vs. unmasked comparison takes the **top-10 aggregated tokens** per profession ($K_2 = 10$) and computes the shared token ratio:

$$r(p) = \frac{|\mathcal{T}_{\text{sh}}|}{|\mathcal{T}_{\text{sh}}| + |\mathcal{T}_{\text{m}}| + |\mathcal{T}_{\text{u}}|} \quad (3)$$

where $\mathcal{T}_{\text{sh}}$, $\mathcal{T}_{\text{m}}$, $\mathcal{T}_{\text{u}}$ are the sets of tokens shared, masked-only, and unmasked-only respectively among the top-10. Gender-conditioned proxy tables are produced identically, stratified by profession and gender field.

Faithfulness validation. The top-5 attributed tokens are erased from each example’s text and the model is re-run to measure the probability drop:

$$\text{Faith}(x) = P(y \mid x) - P(y \mid x_{\text{erased}}) \quad (4)$$

5 Experimental Setup

Code and prediction files are available at <https://github.com/usukhbayarp/comp0087-fairness-occupation-prediction>.

Model	Type	Regime	Cond.
Pythia-160M	Gen.	zero-shot, few-shot, FT	M, U
Pythia-410M	Gen.	zero-shot, few-shot, FT	M, U
Pythia-1.4B	Gen.	zero-shot, few-shot, FT	M, U
DistilBERT	Enc.	Fine-tuned	M, U
RoBERTa	Enc.	Fine-tuned	M, U

Table 3: All model configurations. FT=fine-tuned; M=masked; U=unmasked.

Table 3 summarises all evaluated configurations.

6 Results and Discussion

We evaluate our models across three dimensions: task performance (§6.1), gender fairness (§6.2), and token-level proxy bias mechanisms (§6.3).

The following results are descriptive of only the 3,000-sample evaluation subset, which may introduce some noise.

6.1 Task Performance

Pythia-1.4B fine-tuned (M) reaches $F1 = 0.877$, nearing encoder baselines (RoBERTa (M) 0.880, DistilBERT (M) 0.872); Pythia-160M fine-tuned fails ($F1 \approx 0.02$), likely due to label-space collapse under LoRA at insufficient capacity. Zero-shot $F1$ scales from 0.23 (160M) to 0.57 (1.4B); few-shot

Model	C.	Acc.	F1
Pythia-160M FT	U	0.079	0.020
Pythia-160M FT	M	0.119	0.023
Pythia-410M FT	U	0.581	0.323
Pythia-410M FT	M	0.576	0.325
Pythia-1.4B FT	U	0.877	0.848
Pythia-1.4B FT	M	0.877	0.852
DistilBERT FT	U	0.875	0.843
DistilBERT FT	M	0.872	0.844
RoBERTa FT	U	0.882	0.851
RoBERTa FT	M	0.880	0.851
Pythia-160M zero-shot	U	0.405	0.230
Pythia-160M zero-shot	M	0.400	0.223
Pythia-160M few-shot	U	0.241	0.202
Pythia-160M few-shot	M	0.249	0.206
Pythia-410M zero-shot	U	0.536	0.557
Pythia-410M zero-shot	M	0.533	0.553
Pythia-410M few-shot	U	0.366	0.391
Pythia-410M few-shot	M	0.363	0.389
Pythia-1.4B zero-shot	U	0.658	0.567
Pythia-1.4B zero-shot	M	0.659	0.568
Pythia-1.4B few-shot	U	0.558	0.578
Pythia-1.4B few-shot	M	0.558	0.577

Table 4: Task performance results. U=unmasked; M=masked.

hurts at smaller scales (410M $\rightarrow$ 160M: 0.557 $\rightarrow$ 0.391) but recovers at 1.4B (0.578), suggesting a minimum capacity threshold. Masking has negligible impact on task performance across all models.

6.2 Fairness Results




Figure 1: Pareto frontier of Macro-F1 vs. Equalized Odds gap. RoBERTa masked fine-tuned, followed by DistilBERT, is the most Pareto-optimal model. The generative models which are masked and with largest capacity, as well as the baseline encoders, strongly dominate lower-capacity, zero-shot and few-shot alternatives.

Fine-tuned models achieve markedly lower EO

Model	C.	EO	DP	TPR	FPR
RoBERTa FT	U	0.082	0.024	0.082	0.005
RoBERTa FT	M	0.063	0.022	0.063	0.004
DistilBERT FT	U	0.088	0.022	0.088	0.004
DistilBERT FT	M	0.071	0.022	0.071	0.004
Pythia-1.4B FT	U	0.086	0.023	0.086	0.004
Pythia-1.4B FT	M	0.080	0.022	0.080	0.003
Pythia-1.4B FS	U	0.121	0.031	0.121	0.016
Pythia-1.4B FS	M	0.124	0.031	0.124	0.016
Pythia-1.4B 0s	U	0.210	0.029	0.210	0.019
Pythia-1.4B 0s	M	0.210	0.029	0.210	0.019
Pythia-410m FT	U	0.085	0.012	0.085	0.006
Pythia-410m FT	M	0.088	0.011	0.088	0.006
Pythia-410m FS	U	0.126	0.021	0.126	0.012
Pythia-410m FS	M	0.128	0.020	0.128	0.011
Pythia-410m 0s	U	0.143	0.037	0.143	0.020
Pythia-410m 0s	M	0.131	0.037	0.131	0.021
Pythia-160m FT	U	0.017	0.001	0.017	0.001
Pythia-160m FT	M	0.003	0.002	0.003	0.002
Pythia-160m FS	U	0.080	0.012	0.080	0.007
Pythia-160m FS	M	0.077	0.013	0.077	0.008
Pythia-160m 0s	U	0.092	0.027	0.092	0.025
Pythia-160m 0s	M	0.092	0.027	0.092	0.024

Table 5: Fairness metrics (EO = Equalized Odds gap, DP = Demographic Parity gap, TPR = True Positive Rate gap, FPR = False Positive Rate gap) for each model, both masked (M) and unmasked (U). Row label abbreviations: FT = fine-tuned; FS = few-shot; 0s = zero-shot.




Figure 2: Scaling behavior of Equalized Odds gap across model sizes, 1.4B (red), 410m (green), and 160m (blue). The solid line corresponds to unmasked and dashed to masked. Zero-shot bias increases with scale, while fine-tuning mitigates this effect.

gaps than zero-shot counterparts; zero-shot EO grows from ~ 0.09 (160M) to ~ 0.21 (1.4B), amplifying stereotypes without supervision. Masking reduces aggregate EO for encoders (RoBERTa: 0.082 $\rightarrow$ 0.063; DistilBERT: 0.088 $\rightarrow$ 0.071) but

is non-uniform at profession level: Pythia-1.4B FT sees backfire for *teacher*, *poet*, and *painter*, while Pythia-410M FT shows aggregate backfire (0.085→0.088). FPR gap is consistently much smaller than TPR gap, indicating TPR disparity is the primary fairness issue. To understand why disparities persist under masking, we turn to token-level attribution analysis.

6.3 Proxy Bias Analysis

We analyse which tokens drive occupation predictions and how their importance shifts under gender masking. We first assess attribution stability under masking, then examine which token types carry the gender signal, and finally assess token importance using erasure-based faithfulness analysis.

Profession	DistilBERT	RoBERTa
accountant	0.43	0.82
architect	0.43	0.82
attorney	0.67	0.82
comedian	0.54	0.33
composer	0.54	0.43
dentist	0.54	0.33
dietitian	0.67	0.54
filmmaker	0.67	0.54
journalist	0.67	0.67
model	0.82	0.43
nurse	0.43	0.54
painter	0.67	0.54
photographer	0.67	0.67
physician	0.67	0.54
poet	0.54	0.67
professor	0.82	1.00
psychologist	0.82	0.82
software engineer	0.67	0.43
surgeon	0.67	0.67
teacher	0.82	0.82

Table 6: Shared token ratio (masked vs. unmasked) for all 20 professions. Higher = more stable. Ratio = shared / (shared + masked-only + unmasked-only) over the top-10 aggregated tokens.

Under RoBERTa, *comedian* and *dentist* are the most unstable professions (ratio 0.33), while under DistilBERT both are moderately stable (0.54)—the same reversal observed for *nurse* (DistilBERT 0.43, RoBERTa 0.54). Under DistilBERT, *accountant* and *architect* reach the same minimum ratio (0.43) as *nurse*, while both are substantially more stable under RoBERTa (0.82), illustrating that instability is not confined to gender-skewed professions. *Professor* reaches maximum stability under

RoBERTa (1.00), meaning its top-10 tokens are identical across masked and unmasked conditions. These cross-profession reversals reinforce the conclusion that proxy stability depends on the specific model, not only training data.

When gender tokens are removed, attribution is redistributed across tokens rather than concentrated on the same dominant token, consistent with the completeness property of Integrated Gradients. The profession label token (e.g. *nurse* for the nurse class) frequently appears among the top- K attributed tokens, reflecting label leakage where the profession name occurs directly in the biography. This is expected; the methodologically informative signal lies in how attribution shifts between unmasked and masked conditions. For DistilBERT on *nurse*, the label token *nurse* remains prominent under both conditions (129.89 unmasked, 152.73 masked), consistent with label leakage noted above. The informative shift occurs in secondary tokens: *practitioner* declines from 26.51 to 11.18 and *graduated* from 19.11 to 1.47, while *hospital* (8.96) and *especially* (15.52) gain attribution weight. For RoBERTa on *physician*, masking shifts weight toward higher-level clinical context: *medicine* rises from 166.64 to 292.25 and *patients* (52.72) emerges as a new signal, while *affiliated* (68.78 unmasked) drops from the top-5.

Model	Token	Type	W.Score
DistilBERT	<i>actor</i>	gender	1.15
DistilBERT	<i>sister</i>	gender	1.13
DistilBERT	<i>mrs</i>	gender	1.05
RoBERTa	<i>mrs</i>	gender	2.11
RoBERTa	<i>sister</i>	gender	1.13
DistilBERT	<i>nursing</i>	profession	55.53
DistilBERT	<i>medicine</i>	profession	268.20
DistilBERT	<i>research</i>	profession	329.52
RoBERTa	<i>nursing</i>	profession	70.44
RoBERTa	<i>medicine</i>	profession	166.64
RoBERTa	<i>research</i>	profession	333.32

Table 7: Attribution weight of gender tokens vs. dominant profession tokens, within each model. Gender tokens appear in both models but with substantially smaller weights than profession tokens.

Under both models, a small number of explicit gender tokens appear in the aggregated top- K features, but with substantially smaller weights than dominant profession tokens: within DistilBERT, the highest gender token (*actor*: 1.15) is 48–287×

smaller than DistilBERT profession tokens (55.53–329.52); within RoBERTa, the highest gender token (*mrs*: 2.11) is 33–158× smaller than RoBERTa profession tokens (70.44–333.32). Bias likely operates primarily through indirect contextual signals rather than explicit gender indicators—consistent with an account in which models exploit statistical co-occurrences present in the training corpus [De-Arteaga et al., 2019, Gonen and Goldberg, 2019]: gender-skewed role vocabulary, credential terminology, and contextual framing patterns become spurious correlates that carry the gender signal without any explicit gender reference.

Gender-conditioned analysis reveals partial asymmetry. Under DistilBERT, nurse|F and nurse|M both retain 5 shared tokens (ratio 0.33 each), showing no directional asymmetry. Under RoBERTa, nurse|F retains 7 shared tokens (ratio 0.54) versus nurse|M retaining 6 (ratio 0.43), a modest asymmetry suggesting the learned feature space for *nurse* is somewhat more anchored in female-typical contextual patterns. This attribution-level asymmetry aligns with the fairness results in Section 6.2, where *nurse* shows one of the largest EO gap reductions under masking—suggesting that professions with unstable proxy representations are precisely those where masking provides the most benefit.

Qualitative Illustrations

The following three examples from the DistilBERT unmasked condition provide concrete illustration of the aggregate patterns reported above.

Example 1: Gender asymmetry in attributed tokens (nurse, unmasked). Both biographies are correctly predicted as *nurse*.

Male biography excerpt	Female biography excerpt
“He holds a Masters of Nursing Degree and has served in the Australian Defense force in various roles for over 32 years. He has presented in this field for over 20 years. . . ”	“Her scope of practice encompasses health promotion, disease prevention, diagnosis and management of common and complex health care problems. . . throughout the aging process. . . ”
Top tokens: <i>nursing</i> (3.18), <i>masters</i> (0.61), <i>presented</i> (0.59), <i>australian</i> (0.44)	Top tokens: <i>scope</i> (0.74), <i>health</i> (0.66), <i>aging</i> (0.61), <i>practice</i> (0.55), <i>her</i> (0.49)

The male prediction anchors on explicit credentials (*nursing, masters*); the female prediction dis-

tributes across social-care context words (*scope, aging*)—consistent with the nurse|M faithfulness asymmetry (0.756 vs. 0.354).

Example 2: Prediction flip after masking (backfire).

“[MASK] is a parent of a child with autism, and is dedicated to providing professional assistance to other parents. . . [MASK] can be reached at (408) 843-7997, and is available for paid phone consultation. . . ”

True label: *psychologist*. Unmasked prediction: *psychologist* — top tokens: *parent* (1.15), *autism* (0.81), *child* (0.39), *professional* (0.33). Masked prediction: *attorney* × — top tokens: *408* (0.63), *autism* (0.34), *reached* (0.33), *parents* (0.28).

After masking *He*, the parent/child anchor loses salience; the phone number and “consultation” framing pattern-match *attorney*, producing an error rather than a fairness gain.

Example 3 (mini): Indirect proxy tokens driving a physician prediction.

“Ms. Croteau **practices medicine** in New Britain, CT and specializes in Allergy & Immunology. Ms. Croteau is affiliated with **Hospital** of Central Connecticut.”

Predicted: *physician* — top tokens: *medicine* (1.13), *practices* (0.96), *hospital* (0.52). No gender tokens appear in the top 5. Clinical role words alone drive the prediction, consistent with the interpretation that indirect contextual signals, not explicit gender markers, carry the primary predictive load.

Faithfulness Validation

While Integrated Gradients provides correlational insights into feature importance, the top-*K* erasure test provides intervention-based validation: erasing attributed tokens and measuring the resulting confidence drop assesses whether those tokens are functionally important for the prediction. Table 8 reports mean faithfulness scores by profession across both models and masking conditions. Across all 3,000 examples, 89.6% (DistilBERT) and 82.3% (RoBERTa) of predictions show positive faithfulness, consistent with the attributed tokens being genuinely load-bearing.

Faithfulness scores align with the proxy stability analysis: professions with lower shared ratios tend toward higher faithfulness (*dietitian*: 0.619, *nurse*:

Profession	DistilBERT		RoBERTa	
	U	M	U	M
Dietitian	0.619	0.743	0.548	0.372
Nurse	0.395	0.518	0.203	0.197
Psychologist	0.356	0.378	0.232	0.252
Painter	0.473	0.418	0.339	0.342
Physician	0.233	0.191	0.235	0.200
Attorney	0.162	0.221	0.175	0.158
Professor	0.160	0.214	0.162	0.182
Software Engineer	0.171	0.204	0.193	0.199
<i>Overall</i>	0.330	0.364	0.243	0.249

Table 8: Mean faithfulness (probability drop after erasing top-5 attributed tokens) for a selected subset of professions. U=unmasked; M=masked. Higher = more functionally important tokens under erasure. The *Overall* row is the macro-average across all 20 professions (equal weight per profession, not per example). Full results for all 20 professions are in Appendix E.

0.395), while stable professions show lower faithfulness (*professor*: 0.160, *attorney*: 0.162). This supports the interpretation that attribution instability identifies professions whose predictions are concentrated in a small, functionally decisive feature set. Notably, the *model* profession under RoBERTa unmasked achieves the lowest faithfulness of all 20 professions (0.100), suggesting its predictions are distributed across many weakly-attributing tokens rather than a small decisive set.

Gender-conditioned faithfulness reveals a striking asymmetry for *nurse* under DistilBERT: $\text{nurse}|M = 0.756$ versus $\text{nurse}|F = 0.354$. Erasing the top-5 tokens for male nurse predictions nearly halves the model’s confidence, consistent with male nurse representations being concentrated in a critically small feature set. Under RoBERTa the asymmetry is also present: $\text{nurse}|M = 0.534$ versus $\text{nurse}|F = 0.163$. Full gender-conditioned faithfulness results for all professions with $n \geq 10$ per gender are reported in Appendix D; a second notable asymmetry is observed for *poet* under DistilBERT unmasked (0.434 vs. 0.200).

Masking increases overall faithfulness for both models (DistilBERT: +0.034; RoBERTa: +0.006), suggesting that after gender signals are removed, the remaining attributed tokens become more functionally decisive — consistent with the model concentrating predictions on a smaller, more critical feature set.

7 Conclusion

Fine-tuned models substantially outperform zero/few-shot baselines in both performance and fairness; scaling without supervision amplifies bias; masking reduces aggregate EO gaps in encoder models (by 0.017–0.019) but produces profession-level backfire effects in Pythia; and bias is encoded primarily through indirect contextual tokens whose importance is model-dependent and whose functional load is supported by erasure faithfulness testing. Notably, zero-shot EO gap grows from ~ 0.09 at 160M to ~ 0.21 at 1.4B parameters—suggesting that model capacity alone is not a substitute for task-specific fine-tuning or explicit debiasing interventions. Gender-conditioned analysis further reveals that male nurse predictions are disproportionately dependent on a critically small feature set, with faithfulness nearly double that of female nurse predictions (0.756 vs. 0.354 under DistilBERT).

Limitations. The dataset provides binary gender labels inferred from pronoun usage by the original authors [De-Arteaga et al., 2019], which do not capture the full spectrum of gender identity and limit fairness evaluation to a male/female binary. Attribution is restricted to encoder models. Integrated Gradients and token erasure provide evidence of functional token importance but do not establish causal relationships between input features and model predictions; the scores reflect learned associations, not mechanisms. Per-profession evaluation uses proportional stratification from the 3,000-sample canonical subset, resulting in small sample sizes for low-frequency professions (e.g. comedian $n = 27$, dietitian $n = 30$); per-profession fairness estimates for these occupations should be interpreted with caution. The few-shot exemplars were not gender-balanced (2 male, 1 female pronouns); no consistent directional effect on TPR gaps is observable for the three exemplar professions across model sizes (§4.1, footnote), but a residual exemplar-level confound cannot be fully excluded. Differences in masked/unmasked fairness metrics are reported descriptively without statistical significance tests; small per-profession sample sizes mean some observed gaps may be within sampling noise.

Future work. Remaining directions include extending attribution to generative models to provide mechanistic explanations for observed fairness dis-

parities; adding bootstrapped confidence intervals on fairness metrics to assess the significance of small masked/unmasked differences; and designing a formal cross-gender proxy bias score that jointly captures attribution asymmetry and fairness gaps across professions.

Acknowledgements

We thank our teaching assistant Karen Hambarzumyan for his guidance and support throughout the semester.

A Per-Occupation Proxy Word Tables

Top-5 attributed tokens (weighted score = count_topk $\times$ mean attribution) per profession for DistilBERT and RoBERTa under unmasked (U) and masked (M) conditions. Aggregated over 3,000 test examples; stopwords and tokens shorter than 3 characters excluded. *Note: occasional RoBERTa entries may reflect corpus artefacts (e.g. proper nouns or URL fragments from the training data) rather than semantically meaningful profession features.*

Accountant

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
accounting	20.8	accounting	21.2
financial	9.5	financial	6.7
accountants	4.3	accountants	5.3
accounts	3.8	business	5.3
finance	3.1	tax	5.1
<i>Masked</i>			
accounting	25.2	accounting	30.7
financial	7.3	financial	10.5
accountants	4.6	finance	10.2
finance	2.8	business	7.4
accounts	2.7	cfo	6.2

Architect

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
architecture	29.7	architecture	38.9
architectural	15.5	architectural	24.6
design	13.1	design	17.7
architect	11.1	architect	10.6
architects	7.9	architects	10.4
<i>Masked</i>			
architecture	38.8	architecture	57.1
architectural	19.8	architectural	29.9
architect	12.4	design	24.1
design	10.3	architect	12.5
architects	6.2	architects	12.2

Attorney

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
law	56.0	law	143.6
attorney	37.4	litigation	54.0
legal	27.6	legal	49.5
litigation	25.5	practice	47.1
clients	19.8	attorney	46.2
<i>Masked</i>			
law	82.0	law	182.2
attorney	33.2	litigation	70.7
litigation	30.2	attorney	56.7
legal	25.5	practice	52.3
clients	18.4	legal	31.7

Comedian

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
comedy	16.4	comedy	12.6
stand	4.3	comedians	4.6
comedians	3.0	stand-up	4.4
performed	2.0	joke	3.6
april	1.3	comedic	3.2
<i>Masked</i>			
comedy	16.5	comedy	23.7
stand	4.5	bajwa	8.4
comedians	4.1	comedians	5.0
performed	2.3	stand-up	4.9
jokes	1.9	show	3.4

Composer

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
music	24.3	short	25.5
composition	6.9	music	22.4
composer	6.4	composition	9.6
musical	4.6	compositions	7.9
composing	4.3	musical	7.1
<i>Masked</i>			
music	37.4	music	53.2
composition	9.1	composition	12.2
composer	5.8	compositions	10.9
compositions	5.8	musical	9.5
musical	5.7	orchestra	8.6

Dentist

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
dental	107.2	dental	185.2
dentistry	48.8	dentistry	39.8
dentist	28.9	clinic	25.2
bds	23.4	dentist	24.1
clinic	9.6	practices	19.6
<i>Masked</i>			
dental	107.4	dental	200.2
dentistry	41.1	dentistry	45.8
dentist	32.2	clinic	29.6
patients	9.9	patients	27.9
dentists	8.8	dentist	27.4

Dietitian

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
<i>nutrition</i>	37.7	<i>nutrition</i>	30.7
<i>dietitian</i>	35.3	<i>dietitian</i>	27.8
<i>registered</i>	6.4	<i>professional</i>	10.4
<i>nutritional</i>	3.8	<i>nutritional</i>	5.1
<i>diet</i>	3.2	<i>registered</i>	4.7
<i>Masked</i>			
<i>dietitian</i>	37.5	<i>nutrition</i>	45.7
<i>nutrition</i>	35.7	<i>professional</i>	21.6
<i>registered</i>	7.4	<i>dietitian</i>	14.5
<i>diets</i>	4.8	<i>nutritional</i>	5.0
<i>diet</i>	4.6	<i>appointment</i>	3.8

Filmmaker

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
<i>film</i>	21.0	<i>films</i>	19.7
<i>films</i>	12.5	<i>film</i>	17.1
<i>documentary</i>	8.5	<i>filmmaker</i>	7.8
<i>filmmaking</i>	7.8	<i>feature</i>	7.6
<i>directed</i>	4.6	<i>directed</i>	6.0
<i>Masked</i>			
<i>film</i>	32.2	<i>film</i>	39.5
<i>films</i>	16.4	<i>films</i>	28.8
<i>documentary</i>	7.5	<i>filmmaking</i>	12.0
<i>filmmaking</i>	5.8	<i>filmmaker</i>	9.0
<i>directed</i>	5.5	<i>feature</i>	8.3

Journalist

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
<i>journalism</i>	34.6	<i>journalism</i>	28.7
<i>journalist</i>	18.1	<i>news</i>	16.2
<i>news</i>	13.7	<i>journalist</i>	15.4
<i>editor</i>	11.5	<i>reporting</i>	8.8
<i>journalists</i>	8.6	<i>times</i>	8.4
<i>Masked</i>			
<i>journalism</i>	27.8	<i>journalism</i>	44.6
<i>journalist</i>	16.6	<i>news</i>	24.6
<i>editor</i>	12.8	<i>journalist</i>	22.2
<i>news</i>	10.6	<i>journalists</i>	15.6
<i>reporting</i>	8.6	<i>writing</i>	10.4

Model

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
<i>model</i>	17.1	<i>model</i>	24.2
<i>freeones</i>	15.0	<i>modeling</i>	11.5
<i>modeling</i>	14.2	<i>ranked</i>	10.4
<i>born</i>	7.2	<i>born</i>	8.4
<i>gallery</i>	5.3	<i>modelling</i>	6.0
<i>Masked</i>			
<i>model</i>	12.1	<i>model</i>	31.1
<i>modeling</i>	11.5	<i>born</i>	22.7
<i>born</i>	7.7	<i>gallery</i>	18.1
<i>freeones</i>	7.5	<i>modeling</i>	17.1
<i>ranked</i>	6.0	<i>modelling</i>	8.2

Nurse

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
<i>nurse</i>	129.9	<i>nurse</i>	141.1
<i>nursing</i>	55.5	<i>nursing</i>	70.4
<i>practitioner</i>	26.5	<i>practitioner</i>	25.3
<i>graduated</i>	19.1	<i>graduated</i>	24.9
<i>experiences</i>	15.2	<i>medical</i>	13.3
<i>Masked</i>			
<i>nurse</i>	152.7	<i>nurse</i>	166.4
<i>nursing</i>	60.7	<i>nursing</i>	85.1
<i>especially</i>	15.5	<i>practitioner</i>	20.5
<i>practitioner</i>	11.2	<i>call</i>	16.0
<i>hospital</i>	9.0	<i>medical</i>	14.2

Painter

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
<i>paintings</i>	25.6	<i>painting</i>	36.1
<i>painting</i>	21.0	<i>paintings</i>	24.3
<i>gallery</i>	8.2	<i>art</i>	8.1
<i>art</i>	8.0	<i>gallery</i>	5.4
<i>paint</i>	4.1	<i>paint</i>	4.8
<i>Masked</i>			
<i>paintings</i>	25.7	<i>painting</i>	51.4
<i>painting</i>	24.5	<i>paintings</i>	50.0
<i>art</i>	8.1	<i>art</i>	13.3
<i>gallery</i>	4.7	<i>gallery</i>	9.6
<i>work</i>	4.4	<i>depicted</i>	8.7

Photographer

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
<i>photography</i>	149.1	<i>photography</i>	126.0
<i>photographs</i>	26.2	<i>photographs</i>	23.9
<i>photographer</i>	23.6	<i>photographer</i>	23.0
<i>photographic</i>	16.2	<i>photographic</i>	15.1
<i>photo</i>	14.3	<i>shoots</i>	12.8
<i>Masked</i>			
<i>photography</i>	143.0	<i>photography</i>	239.6
<i>photographs</i>	24.3	<i>photographs</i>	45.2
<i>photographer</i>	22.0	<i>photographer</i>	39.2
<i>work</i>	15.9	<i>photographic</i>	29.4
<i>photographic</i>	15.8	<i>images</i>	24.9

Physician

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
<i>medicine</i>	268.2	<i>practices</i>	179.3
<i>practices</i>	76.9	<i>medicine</i>	166.6
<i>medical</i>	66.4	<i>medical</i>	77.4
<i>physician</i>	36.1	<i>affiliated</i>	68.8
<i>graduate</i>	15.9	<i>physician</i>	26.0
<i>Masked</i>			
<i>medicine</i>	287.6	<i>medicine</i>	292.2
<i>medical</i>	89.3	<i>practices</i>	224.6
<i>practices</i>	72.8	<i>medical</i>	97.4
<i>hospital</i>	33.0	<i>patients</i>	52.7
<i>physician</i>	31.0	<i>specializes</i>	38.3

Poet

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
<i>poetry</i>	34.9	<i>poetry</i>	52.0
<i>poems</i>	8.1	<i>poems</i>	15.7
<i>writing</i>	4.7	<i>writing</i>	8.7
<i>english</i>	3.1	<i>poem</i>	6.3
<i>poem</i>	3.0	<i>published</i>	3.6
<i>Masked</i>			
<i>poetry</i>	44.9	<i>poetry</i>	67.8
<i>poems</i>	9.9	<i>poems</i>	14.0
<i>writing</i>	5.1	<i>writing</i>	7.6
<i>poem</i>	4.5	<i>poem</i>	6.6
<i>english</i>	3.5	<i>review</i>	6.6

Professor

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
<i>research</i>	329.5	<i>research</i>	333.3
<i>phd</i>	93.8	<i>ph.d</i>	170.5
<i>interests</i>	60.9	<i>phd</i>	119.2
<i>university</i>	58.6	<i>university</i>	105.0
<i>focuses</i>	49.3	<i>interests</i>	48.4
<i>Masked</i>			
<i>research</i>	377.2	<i>research</i>	514.2
<i>university</i>	193.5	<i>university</i>	302.8
<i>phd</i>	79.8	<i>interests</i>	95.4
<i>professor</i>	49.0	<i>phd</i>	94.1
<i>interests</i>	46.9	<i>ph.d</i>	91.7

Psychologist

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
<i>psychologist</i>	58.6	<i>psychologist</i>	55.0
<i>psychology</i>	43.0	<i>psychology</i>	31.5
<i>psychological</i>	9.2	<i>clinical</i>	18.7
<i>psychotherapy</i>	9.2	<i>appointment</i>	13.2
<i>experiences</i>	8.3	<i>experiences</i>	10.5
<i>Masked</i>			
<i>psychologist</i>	56.4	<i>psychologist</i>	61.6
<i>psychology</i>	33.0	<i>psychology</i>	40.4
<i>clinical</i>	17.5	<i>clinical</i>	22.0
<i>psychological</i>	8.4	<i>psychological</i>	16.8
<i>psychotherapy</i>	7.1	<i>psychotherapy</i>	13.4

Software Engineer

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
<i>software</i>	4.8	<i>development</i>	12.3
<i>development</i>	4.6	<i>engineering</i>	9.7
<i>engineering</i>	4.2	<i>computer</i>	6.4
<i>ibm</i>	2.9	<i>software</i>	6.3
<i>code</i>	2.6	<i>building</i>	3.4
<i>Masked</i>			
<i>computer</i>	7.4	<i>software</i>	14.2
<i>engineering</i>	7.4	<i>engineering</i>	8.0
<i>software</i>	5.8	<i>development</i>	4.7
<i>development</i>	4.5	<i>javascript</i>	4.5
<i>intel</i>	3.7	<i>devops</i>	4.4

Surgeon

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
<i>surgery</i>	44.2	<i>surgery</i>	51.1
<i>medical</i>	13.2	<i>medical</i>	20.0
<i>medicine</i>	12.1	<i>practice</i>	16.5
<i>practices</i>	8.9	<i>medicine</i>	12.2
<i>surgeon</i>	7.7	<i>hospital</i>	10.8
<i>Masked</i>			
<i>surgery</i>	24.4	<i>surgery</i>	95.6
<i>medical</i>	15.9	<i>medical</i>	22.9
<i>practice</i>	12.7	<i>surgeon</i>	16.8
<i>medicine</i>	11.7	<i>medicine</i>	13.9
<i>orthopaedic</i>	8.0	<i>neurosurgery</i>	10.6

Teacher

DistilBERT (U/M)		RoBERTa (U/M)	
Token	W.Score	Token	W.Score
<i>Unmasked</i>			
<i>teacher</i>	19.9	<i>students</i>	25.3
<i>students</i>	18.7	<i>teaching</i>	20.7
<i>teaching</i>	16.3	<i>teacher</i>	20.4
<i>school</i>	12.8	<i>school</i>	18.6
<i>teachers</i>	12.4	<i>taught</i>	13.7
<i>Masked</i>			
<i>students</i>	27.1	<i>students</i>	38.4
<i>teacher</i>	18.2	<i>teaching</i>	34.2
<i>teaching</i>	14.5	<i>school</i>	26.3
<i>teachers</i>	12.3	<i>teaches</i>	26.1
<i>school</i>	10.8	<i>teacher</i>	23.6

B Prompt Templates

B.1 Zero-Shot

"""

Given the short biography below, predict the occupation of the person. Answer with a single specific occupation word (e.g. professor, physician, attorney, photographer, journalist, nurse, psychologist, teacher, dentist, surgeon, architect, painter, model, poet, filmmaker, software_engineer, accountant, composer, dietitian, comedian etc.).

Biography: {text}

Occupation:

"""

B.2 Few-Shot

"""

Given the short biography below, predict the occupation of the person. Answer with a single specific occupation word (e.g. professor, physician, attorney, photographer, journalist, nurse, psychologist, teacher, dentist, surgeon, architect, painter, model, poet, filmmaker, software_engineer, accountant, composer, dietitian, comedian etc.).

Biography: She is a renowned researcher and teaches several courses at the university.

Occupation: professor

Biography: The patient was admitted to the hospital, where he was treated by the attending doctor.

Occupation: physician

Biography: He covers local news and politics for the daily newspaper.

Occupation: journalist

Biography: {text}

Occupation:

"""

C Detailed Fairness Metrics by Occupation

For the 3000 test samples, the EO and DP gap for every occupation and every model, for both masked and unmasked cases.

Accountant

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.094	0.002	0.012	Pythia-410m FS	U	0.167	0.002	0.009
DistilBERT FT	M	0.031	0.000	0.009	Pythia-410m FS	M	0.167	0.002	0.009
RoBERTa FT	U	0.062	0.000	0.009	Pythia-410m 0s	U	0.073	0.001	0.004
RoBERTa FT	M	0.010	0.001	0.008	Pythia-410m 0s	M	0.073	0.000	0.003
Pythia-1.4B FT	U	0.229	0.002	0.013	Pythia-160m FT	U	0.000	0.000	0.000
Pythia-1.4B FT	M	0.229	0.002	0.013	Pythia-160m FT	M	0.000	0.000	0.000
Pythia-1.4B FS	U	0.333	0.004	0.015	Pythia-160m FS	U	0.083	0.000	0.001
Pythia-1.4B FS	M	0.333	0.005	0.016	Pythia-160m FS	M	0.083	0.000	0.001
Pythia-1.4B 0s	U	0.260	0.001	0.001	Pythia-160m 0s	U	0.000	0.000	0.000
Pythia-1.4B 0s	M	0.260	0.001	0.001	Pythia-160m 0s	M	0.000	0.000	0.000
Pythia-410m FT	U	0.062	0.003	0.001					
Pythia-410m FT	M	0.042	0.001	0.003					

Architect

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.242	0.006	0.017	Pythia-410m FS	U	0.118	0.010	0.021
DistilBERT FT	M	0.134	0.008	0.019	Pythia-410m FS	M	0.170	0.010	0.022
RoBERTa FT	U	0.173	0.008	0.020	Pythia-410m 0s	U	0.082	0.059	0.076
RoBERTa FT	M	0.156	0.007	0.020	Pythia-410m 0s	M	0.099	0.060	0.077
Pythia-1.4B FT	U	0.209	0.005	0.017	Pythia-160m FT	U	0.000	0.000	0.000
Pythia-1.4B FT	M	0.175	0.003	0.016	Pythia-160m FT	M	0.000	0.002	0.002
Pythia-1.4B FS	U	0.012	0.011	0.021	Pythia-160m FS	U	0.131	0.002	0.008
Pythia-1.4B FS	M	0.012	0.010	0.020	Pythia-160m FS	M	0.114	0.001	0.006
Pythia-1.4B 0s	U	0.135	0.002	0.005	Pythia-160m 0s	U	0.014	0.001	0.002
Pythia-1.4B 0s	M	0.135	0.002	0.005	Pythia-160m 0s	M	0.050	0.001	0.003
Pythia-410m FT	U	0.182	0.015	0.015					
Pythia-410m FT	M	0.028	0.002	0.005					

Attorney

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.046	0.001	0.019	Pythia-410m FS	U	0.153	0.004	0.018
DistilBERT FT	M	0.002	0.000	0.016	Pythia-410m FS	M	0.147	0.004	0.017
RoBERTa FT	U	0.026	0.000	0.018	Pythia-410m 0s	U	0.202	0.008	0.037
RoBERTa FT	M	0.030	0.003	0.016	Pythia-410m 0s	M	0.185	0.009	0.036
Pythia-1.4B FT	U	0.014	0.004	0.014	Pythia-160m FT	U	0.022	0.000	0.003
Pythia-1.4B FT	M	0.020	0.001	0.017	Pythia-160m FT	M	0.002	0.002	0.002
Pythia-1.4B FS	U	0.018	0.008	0.007	Pythia-160m FS	U	0.053	0.001	0.004
Pythia-1.4B FS	M	0.018	0.008	0.007	Pythia-160m FS	M	0.046	0.001	0.003
Pythia-1.4B 0s	U	0.065	0.005	0.001	Pythia-160m 0s	U	0.005	0.000	0.002
Pythia-1.4B 0s	M	0.065	0.005	0.001	Pythia-160m 0s	M	0.012	0.001	0.003
Pythia-410m FT	U	0.009	0.005	0.015					
Pythia-410m FT	M	0.042	0.004	0.010					

Comedian

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.059	0.004	0.010	Pythia-410m FS	U	0.271	0.002	0.006
DistilBERT FT	M	0.118	0.002	0.008	Pythia-410m FS	M	0.271	0.001	0.005
RoBERTa FT	U	0.141	0.003	0.010	Pythia-410m 0s	U	0.271	0.001	0.005
RoBERTa FT	M	0.141	0.002	0.009	Pythia-410m 0s	M	0.271	0.001	0.005
Pythia-1.4B FT	U	0.141	0.001	0.008	Pythia-160m FT	U	0.000	0.000	0.000
Pythia-1.4B FT	M	0.141	0.000	0.007	Pythia-160m FT	M	0.000	0.000	0.000
Pythia-1.4B FS	U	0.365	0.001	0.007	Pythia-160m FS	U	0.200	0.001	0.000
Pythia-1.4B FS	M	0.365	0.001	0.007	Pythia-160m FS	M	0.200	0.000	0.001
Pythia-1.4B 0s	U	0.047	0.000	0.002	Pythia-160m 0s	U	0.235	0.000	0.002
Pythia-1.4B 0s	M	0.047	0.000	0.002	Pythia-160m 0s	M	0.235	0.000	0.002
Pythia-410m FT	U	0.000	0.002	0.002					
Pythia-410m FT	M	0.000	0.002	0.002					

Composer

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.025	0.001	0.015	Pythia-410m FS	U	0.152	0.003	0.011
DistilBERT FT	M	0.025	0.001	0.015	Pythia-410m FS	M	0.152	0.005	0.013
RoBERTa FT	U	0.003	0.000	0.013	Pythia-410m 0s	U	0.003	0.002	0.016
RoBERTa FT	M	0.054	0.001	0.014	Pythia-410m 0s	M	0.003	0.003	0.016
Pythia-1.4B FT	U	0.025	0.001	0.015	Pythia-160m FT	U	0.000	0.000	0.000
Pythia-1.4B FT	M	0.054	0.000	0.015	Pythia-160m FT	M	0.000	0.000	0.000
Pythia-1.4B FS	U	0.162	0.006	0.019	Pythia-160m FS	U	0.051	0.001	0.003
Pythia-1.4B FS	M	0.162	0.006	0.019	Pythia-160m FS	M	0.051	0.001	0.003
Pythia-1.4B 0s	U	0.067	0.003	0.011	Pythia-160m 0s	U	0.051	0.001	0.002
Pythia-1.4B 0s	M	0.067	0.003	0.011	Pythia-160m 0s	M	0.022	0.001	0.002
Pythia-410m FT	U	0.165	0.004	0.004					
Pythia-410m FT	M	0.025	0.001	0.002					

Dentist

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.057	0.002	0.030	Pythia-410m FS	U	0.016	0.000	0.026
DistilBERT FT	M	0.080	0.002	0.029	Pythia-410m FS	M	0.016	0.001	0.027
RoBERTa FT	U	0.068	0.002	0.029	Pythia-410m 0s	U	0.006	0.002	0.026
RoBERTa FT	M	0.057	0.002	0.030	Pythia-410m 0s	M	0.006	0.002	0.025
Pythia-1.4B FT	U	0.068	0.002	0.029	Pythia-160m FT	U	0.000	0.000	0.000
Pythia-1.4B FT	M	0.068	0.002	0.029	Pythia-160m FT	M	0.000	0.000	0.000
Pythia-1.4B FS	U	0.045	0.004	0.031	Pythia-160m FS	U	0.201	0.001	0.021
Pythia-1.4B FS	M	0.045	0.004	0.030	Pythia-160m FS	M	0.198	0.001	0.024
Pythia-1.4B 0s	U	0.019	0.001	0.028	Pythia-160m 0s	U	0.016	0.000	0.016
Pythia-1.4B 0s	M	0.019	0.000	0.027	Pythia-160m 0s	M	0.027	0.001	0.016
Pythia-410m FT	U	0.082	0.002	0.019					
Pythia-410m FT	M	0.053	0.001	0.020					

Dietitian

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.103	0.000	0.018	Pythia-410m FS	U	0.552	0.002	0.010
DistilBERT FT	M	0.034	0.000	0.019	Pythia-410m FS	M	0.552	0.002	0.010
RoBERTa FT	U	0.069	0.001	0.019	Pythia-410m 0s	U	0.586	0.001	0.013
RoBERTa FT	M	0.103	0.000	0.018	Pythia-410m 0s	M	0.552	0.001	0.012
Pythia-1.4B FT	U	0.069	0.001	0.019	Pythia-160m FT	U	0.190	0.002	0.003
Pythia-1.4B FT	M	0.069	0.000	0.018	Pythia-160m FT	M	0.017	0.011	0.013
Pythia-1.4B FS	U	0.103	0.010	0.028	Pythia-160m FS	U	0.000	0.003	0.003
Pythia-1.4B FS	M	0.103	0.010	0.028	Pythia-160m FS	M	0.000	0.004	0.004
Pythia-1.4B 0s	U	0.862	0.000	0.018	Pythia-160m 0s	U	0.000	0.001	0.001
Pythia-1.4B 0s	M	0.862	0.000	0.018	Pythia-160m 0s	M	0.000	0.001	0.001
Pythia-410m FT	U	0.034	0.006	0.007					
Pythia-410m FT	M	0.034	0.003	0.004					

Filmmaker

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.079	0.003	0.007	Pythia-410m FS	U	0.071	0.003	0.002
DistilBERT FT	M	0.021	0.003	0.009	Pythia-410m FS	M	0.071	0.003	0.002
RoBERTa FT	U	0.064	0.003	0.008	Pythia-410m 0s	U	0.236	0.001	0.006
RoBERTa FT	M	0.043	0.003	0.009	Pythia-410m 0s	M	0.157	0.001	0.005
Pythia-1.4B FT	U	0.036	0.001	0.008	Pythia-160m FT	U	0.000	0.001	0.001
Pythia-1.4B FT	M	0.014	0.002	0.009	Pythia-160m FT	M	0.000	0.001	0.001
Pythia-1.4B FS	U	0.093	0.001	0.004	Pythia-160m FS	U	0.114	0.000	0.001
Pythia-1.4B FS	M	0.143	0.002	0.004	Pythia-160m FS	M	0.086	0.000	0.001
Pythia-1.4B 0s	U	0.071	0.000	0.001	Pythia-160m 0s	U	0.029	0.000	0.001
Pythia-1.4B 0s	M	0.071	0.000	0.001	Pythia-160m 0s	M	0.000	0.000	0.000
Pythia-410m FT	U	0.029	0.003	0.003					
Pythia-410m FT	M	0.029	0.002	0.002					

Journalist

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.076	0.003	0.012	Pythia-410m FS	U	0.106	0.037	0.047
DistilBERT FT	M	0.063	0.002	0.013	Pythia-410m FS	M	0.092	0.025	0.035
RoBERTa FT	U	0.111	0.001	0.019	Pythia-410m 0s	U	0.025	0.018	0.025
RoBERTa FT	M	0.048	0.004	0.010	Pythia-410m 0s	M	0.047	0.021	0.027
Pythia-1.4B FT	U	0.055	0.002	0.013	Pythia-160m FT	U	0.000	0.000	0.000
Pythia-1.4B FT	M	0.033	0.007	0.006	Pythia-160m FT	M	0.000	0.000	0.000
Pythia-1.4B FS	U	0.103	0.036	0.030	Pythia-160m FS	U	0.021	0.075	0.077
Pythia-1.4B FS	M	0.090	0.034	0.028	Pythia-160m FS	M	0.007	0.085	0.086
Pythia-1.4B 0s	U	0.350	0.004	0.028	Pythia-160m 0s	U	0.137	0.002	0.011
Pythia-1.4B 0s	M	0.350	0.004	0.028	Pythia-160m 0s	M	0.074	0.002	0.008
Pythia-410m FT	U	0.009	0.002	0.003					
Pythia-410m FT	M	0.043	0.006	0.002					

Model

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.173	0.002	0.028	Pythia-410m FS	U	0.127	0.014	0.001
DistilBERT FT	M	0.304	0.000	0.028	Pythia-410m FS	M	0.127	0.014	0.001
RoBERTa FT	U	0.153	0.002	0.023	Pythia-410m 0s	U	0.227	0.001	0.011
RoBERTa FT	M	0.062	0.001	0.027	Pythia-410m 0s	M	0.227	0.001	0.011
Pythia-1.4B FT	U	0.173	0.001	0.027	Pythia-160m FT	U	0.000	0.001	0.001
Pythia-1.4B FT	M	0.082	0.001	0.027	Pythia-160m FT	M	0.000	0.001	0.001
Pythia-1.4B FS	U	0.233	0.004	0.031	Pythia-160m FS	U	0.069	0.016	0.012
Pythia-1.4B FS	M	0.253	0.005	0.032	Pythia-160m FS	M	0.049	0.016	0.013
Pythia-1.4B 0s	U	0.067	0.004	0.009	Pythia-160m 0s	U	0.162	0.002	0.002
Pythia-1.4B 0s	M	0.067	0.004	0.009	Pythia-160m 0s	M	0.031	0.002	0.000
Pythia-410m FT	U	0.209	0.003	0.008					
Pythia-410m FT	M	0.380	0.003	0.010					

Nurse

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.060	0.012	0.086	Pythia-410m FS	U	0.179	0.029	0.085
DistilBERT FT	M	0.074	0.006	0.078	Pythia-410m FS	M	0.235	0.028	0.085
RoBERTa FT	U	0.067	0.009	0.084	Pythia-410m 0s	U	0.056	0.020	0.084
RoBERTa FT	M	0.019	0.011	0.082	Pythia-410m 0s	M	0.034	0.021	0.083
Pythia-1.4B FT	U	0.100	0.012	0.086	Pythia-160m FT	U	0.056	0.001	0.000
Pythia-1.4B FT	M	0.034	0.009	0.079	Pythia-160m FT	M	0.033	0.001	0.000
Pythia-1.4B FS	U	0.024	0.062	0.123	Pythia-160m FS	U	0.137	0.004	0.017
Pythia-1.4B FS	M	0.024	0.062	0.123	Pythia-160m FS	M	0.107	0.006	0.015
Pythia-1.4B 0s	U	0.660	0.014	0.085	Pythia-160m 0s	U	0.260	0.006	0.030
Pythia-1.4B 0s	M	0.660	0.015	0.086	Pythia-160m 0s	M	0.260	0.007	0.031
Pythia-410m FT	U	0.113	0.005	0.040					
Pythia-410m FT	M	0.232	0.004	0.037					

Painter

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.178	0.001	0.005	Pythia-410m FS	U	0.044	0.003	0.004
DistilBERT FT	M	0.046	0.002	0.000	Pythia-410m FS	M	0.025	0.003	0.002
RoBERTa FT	U	0.057	0.003	0.002	Pythia-410m 0s	U	0.030	0.000	0.001
RoBERTa FT	M	0.046	0.001	0.000	Pythia-410m 0s	M	0.004	0.001	0.000
Pythia-1.4B FT	U	0.043	0.000	0.002	Pythia-160m FT	U	0.000	0.001	0.001
Pythia-1.4B FT	M	0.078	0.001	0.001	Pythia-160m FT	M	0.000	0.000	0.000
Pythia-1.4B FS	U	0.001	0.000	0.000	Pythia-160m FS	U	0.018	0.001	0.001
Pythia-1.4B FS	M	0.001	0.000	0.000	Pythia-160m FS	M	0.053	0.002	0.001
Pythia-1.4B 0s	U	0.105	0.003	0.000	Pythia-160m 0s	U	0.065	0.017	0.015
Pythia-1.4B 0s	M	0.139	0.003	0.001	Pythia-160m 0s	M	0.096	0.020	0.017
Pythia-410m FT	U	0.059	0.003	0.004					
Pythia-410m FT	M	0.018	0.004	0.004					

Photographer

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.013	0.001	0.023	Pythia-410m FS	U	0.079	0.044	0.061
DistilBERT FT	M	0.001	0.002	0.022	Pythia-410m FS	M	0.065	0.038	0.055
RoBERTa FT	U	0.054	0.002	0.029	Pythia-410m 0s	U	0.040	0.004	0.026
RoBERTa FT	M	0.030	0.002	0.023	Pythia-410m 0s	M	0.031	0.003	0.024
Pythia-1.4B FT	U	0.065	0.000	0.026	Pythia-160m FT	U	0.014	0.002	0.003
Pythia-1.4B FT	M	0.028	0.000	0.025	Pythia-160m FT	M	0.000	0.000	0.000
Pythia-1.4B FS	U	0.019	0.014	0.033	Pythia-160m FS	U	0.084	0.003	0.013
Pythia-1.4B FS	M	0.033	0.015	0.035	Pythia-160m FS	M	0.123	0.001	0.018
Pythia-1.4B 0s	U	0.057	0.008	0.008	Pythia-160m 0s	U	0.062	0.000	0.003
Pythia-1.4B 0s	M	0.057	0.008	0.008	Pythia-160m 0s	M	0.035	0.001	0.004
Pythia-410m FT	U	0.046	0.005	0.024					
Pythia-410m FT	M	0.086	0.022	0.041					

Physician

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.087	0.006	0.027	Pythia-410m FS	U	0.055	0.033	0.020
DistilBERT FT	M	0.113	0.002	0.032	Pythia-410m FS	M	0.034	0.030	0.016
RoBERTa FT	U	0.087	0.001	0.031	Pythia-410m 0s	U	0.147	0.043	0.036
RoBERTa FT	M	0.076	0.003	0.028	Pythia-410m 0s	M	0.142	0.052	0.043
Pythia-1.4B FT	U	0.071	0.003	0.027	Pythia-160m FT	U	0.000	0.001	0.001
Pythia-1.4B FT	M	0.083	0.002	0.029	Pythia-160m FT	M	0.006	0.003	0.002
Pythia-1.4B FS	U	0.065	0.014	0.011	Pythia-160m FS	U	0.181	0.003	0.010
Pythia-1.4B FS	M	0.059	0.015	0.010	Pythia-160m FS	M	0.153	0.006	0.003
Pythia-1.4B 0s	U	0.164	0.016	0.021	Pythia-160m 0s	U	0.195	0.016	0.003
Pythia-1.4B 0s	M	0.151	0.016	0.020	Pythia-160m 0s	M	0.201	0.016	0.003
Pythia-410m FT	U	0.104	0.013	0.018					
Pythia-410m FT	M	0.129	0.015	0.018					

Poet

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.220	0.000	0.004	Pythia-410m FS	U	0.107	0.002	0.004
DistilBERT FT	M	0.153	0.002	0.004	Pythia-410m FS	M	0.107	0.004	0.006
RoBERTa FT	U	0.153	0.001	0.003	Pythia-410m 0s	U	0.147	0.000	0.003
RoBERTa FT	M	0.120	0.002	0.004	Pythia-410m 0s	M	0.113	0.001	0.001
Pythia-1.4B FT	U	0.120	0.001	0.003	Pythia-160m FT	U	0.000	0.000	0.000
Pythia-1.4B FT	M	0.187	0.002	0.005	Pythia-160m FT	M	0.000	0.000	0.000
Pythia-1.4B FS	U	0.153	0.001	0.003	Pythia-160m FS	U	0.060	0.003	0.004
Pythia-1.4B FS	M	0.153	0.000	0.003	Pythia-160m FS	M	0.060	0.004	0.005
Pythia-1.4B 0s	U	0.133	0.002	0.004	Pythia-160m 0s	U	0.020	0.008	0.009
Pythia-1.4B 0s	M	0.133	0.003	0.005	Pythia-160m 0s	M	0.093	0.008	0.010
Pythia-410m FT	U	0.060	0.001	0.002					
Pythia-410m FT	M	0.127	0.004	0.002					

Professor

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.024	0.012	0.022	Pythia-410m FS	U	0.040	0.006	0.018
DistilBERT FT	M	0.022	0.024	0.031	Pythia-410m FS	M	0.042	0.005	0.018
RoBERTa FT	U	0.018	0.024	0.032	Pythia-410m 0s	U	0.299	0.029	0.118
RoBERTa FT	M	0.020	0.019	0.028	Pythia-410m 0s	M	0.301	0.028	0.118
Pythia-1.4B FT	U	0.003	0.012	0.030	Pythia-160m FT	U	0.046	0.011	0.005
Pythia-1.4B FT	M	0.003	0.002	0.023	Pythia-160m FT	M	0.006	0.017	0.014
Pythia-1.4B FS	U	0.120	0.010	0.050	Pythia-160m FS	U	0.097	0.019	0.047
Pythia-1.4B FS	M	0.119	0.010	0.050	Pythia-160m FS	M	0.120	0.024	0.057
Pythia-1.4B 0s	U	0.078	0.243	0.206	Pythia-160m 0s	U	0.144	0.231	0.211
Pythia-1.4B 0s	M	0.083	0.242	0.207	Pythia-160m 0s	M	0.165	0.215	0.206
Pythia-410m FT	U	0.020	0.026	0.007					
Pythia-410m FT	M	0.020	0.026	0.008					

Psychologist

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.006	0.002	0.023	Pythia-410m FS	U	0.060	0.001	0.010
DistilBERT FT	M	0.067	0.003	0.021	Pythia-410m FS	M	0.067	0.001	0.010
RoBERTa FT	U	0.034	0.000	0.024	Pythia-410m 0s	U	0.024	0.004	0.014
RoBERTa FT	M	0.002	0.001	0.024	Pythia-410m 0s	M	0.013	0.004	0.015
Pythia-1.4B FT	U	0.060	0.002	0.025	Pythia-160m FT	U	0.011	0.002	0.002
Pythia-1.4B FT	M	0.031	0.002	0.024	Pythia-160m FT	M	0.000	0.001	0.001
Pythia-1.4B FS	U	0.010	0.023	0.038	Pythia-160m FS	U	0.026	0.005	0.005
Pythia-1.4B FS	M	0.010	0.023	0.038	Pythia-160m FS	M	0.008	0.003	0.004
Pythia-1.4B 0s	U	0.305	0.006	0.030	Pythia-160m 0s	U	0.003	0.000	0.002
Pythia-1.4B 0s	M	0.294	0.006	0.029	Pythia-160m 0s	M	0.028	0.000	0.002
Pythia-410m FT	U	0.175	0.011	0.025					
Pythia-410m FT	M	0.120	0.006	0.017					

Software Engineer

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.070	0.010	0.028	Pythia-410m FS	U	0.021	0.003	0.003
DistilBERT FT	M	0.070	0.007	0.025	Pythia-410m FS	M	0.043	0.003	0.004
RoBERTa FT	U	0.112	0.009	0.026	Pythia-410m 0s	U	0.155	0.006	0.013
RoBERTa FT	M	0.134	0.008	0.024	Pythia-410m 0s	M	0.112	0.006	0.012
Pythia-1.4B FT	U	0.027	0.009	0.029	Pythia-160m FT	U	0.000	0.001	0.001
Pythia-1.4B FT	M	0.036	0.007	0.029	Pythia-160m FT	M	0.000	0.001	0.001
Pythia-1.4B FS	U	0.322	0.040	0.061	Pythia-160m FS	U	0.021	0.002	0.003
Pythia-1.4B FS	M	0.322	0.040	0.061	Pythia-160m FS	M	0.021	0.004	0.005
Pythia-1.4B 0s	U	0.103	0.014	0.027	Pythia-160m 0s	U	0.000	0.001	0.001
Pythia-1.4B 0s	M	0.103	0.014	0.027	Pythia-160m 0s	M	0.000	0.001	0.001
Pythia-410m FT	U	0.106	0.001	0.004					
Pythia-410m FT	M	0.043	0.000	0.001					

Surgeon

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.070	0.010	0.041	Pythia-410m FS	U	0.045	0.002	0.010
DistilBERT FT	M	0.020	0.007	0.037	Pythia-410m FS	M	0.007	0.002	0.010
RoBERTa FT	U	0.043	0.011	0.042	Pythia-410m 0s	U	0.018	0.001	0.010
RoBERTa FT	M	0.018	0.007	0.037	Pythia-410m 0s	M	0.018	0.002	0.011
Pythia-1.4B FT	U	0.116	0.012	0.045	Pythia-160m FT	U	0.000	0.000	0.000
Pythia-1.4B FT	M	0.055	0.009	0.041	Pythia-160m FT	M	0.000	0.000	0.000
Pythia-1.4B FS	U	0.002	0.013	0.030	Pythia-160m FS	U	0.016	0.001	0.002
Pythia-1.4B FS	M	0.002	0.013	0.030	Pythia-160m FS	M	0.004	0.001	0.002
Pythia-1.4B 0s	U	0.116	0.001	0.000	Pythia-160m 0s	U	0.000	0.000	0.000
Pythia-1.4B 0s	M	0.104	0.001	0.001	Pythia-160m 0s	M	0.000	0.000	0.000
Pythia-410m FT	U	0.198	0.011	0.030					
Pythia-410m FT	M	0.252	0.005	0.024					

Teacher

Model	C.	TPR	FPR	DP	Model	C.	TPR	FPR	DP
DistilBERT FT	U	0.073	0.000	0.021	Pythia-410m FS	U	0.157	0.033	0.045
DistilBERT FT	M	0.040	0.006	0.025	Pythia-410m FS	M	0.170	0.035	0.047
RoBERTa FT	U	0.132	0.010	0.033	Pythia-410m 0s	U	0.227	0.202	0.220
RoBERTa FT	M	0.086	0.004	0.025	Pythia-410m 0s	M	0.227	0.208	0.225
Pythia-1.4B FT	U	0.099	0.010	0.032	Pythia-160m FT	U	0.000	0.000	0.000
Pythia-1.4B FT	M	0.178	0.009	0.034	Pythia-160m FT	M	0.000	0.000	0.000
Pythia-1.4B FS	U	0.228	0.054	0.074	Pythia-160m FS	U	0.037	0.004	0.008
Pythia-1.4B FS	M	0.228	0.054	0.074	Pythia-160m FS	M	0.056	0.007	0.012
Pythia-1.4B 0s	U	0.534	0.060	0.091	Pythia-160m 0s	U	0.447	0.205	0.225
Pythia-1.4B 0s	M	0.534	0.061	0.092	Pythia-160m 0s	M	0.513	0.205	0.228
Pythia-410m FT	U	0.043	0.001	0.005					
Pythia-410m FT	M	0.053	0.010	0.017					

D Full Faithfulness Scores by Profession

Mean faithfulness (probability drop after erasing top-5 attributed tokens) for all 20 professions across both models and masking conditions. U = unmasked; M = masked. Macro-average (equal weight per profession) reported in the *Overall* row.

Profession	DistilBERT		RoBERTa	
	U	M	U	M
Accountant	0.399	0.422	0.259	0.268
Architect	0.370	0.433	0.283	0.268
Attorney	0.162	0.221	0.175	0.158
Comedian	0.352	0.296	0.240	0.292
Composer	0.373	0.467	0.228	0.329
Dentist	0.353	0.390	0.271	0.228
Dietitian	0.619	0.743	0.548	0.372
Filmmaker	0.277	0.368	0.190	0.188
Journalist	0.268	0.275	0.281	0.318
Model	0.294	0.153	0.100	0.192
Nurse	0.395	0.518	0.203	0.197
Painter	0.473	0.418	0.339	0.342
Photographer	0.385	0.469	0.283	0.346
Physician	0.233	0.191	0.235	0.200
Poet	0.315	0.368	0.124	0.150
Professor	0.160	0.214	0.162	0.182
Psychologist	0.356	0.378	0.232	0.252
Software Engineer	0.171	0.204	0.193	0.199
Surgeon	0.255	0.411	0.249	0.221
Teacher	0.383	0.342	0.269	0.276
<i>Overall</i>	0.330	0.364	0.243	0.249

E Gender-Conditioned Faithfulness by Profession

Mean faithfulness scores split by gender for professions with at least 10 examples per gender group, shown for all four conditions (two models $\times$ masked/unmasked). n = number of examples. Dietitian is excluded from masked rows (male $n=3$ in both masked conditions).

Profession	Female		Male	
	Faith.	n	Faith.	n
<i>DistilBERT — unmasked</i>				
Accountant	0.373	10	0.408	31
Architect	0.490	17	0.326	47
Attorney	0.161	110	0.162	160
Dentist	0.350	31	0.354	85
Dietitian	0.656	27	0.288	3
Filmmaker	0.263	17	0.284	31
Journalist	0.238	80	0.301	74
Nurse	0.354	132	0.756	15
Painter	0.492	31	0.452	28
Photographer	0.432	76	0.357	126
Physician	0.198	171	0.270	157
Poet	0.200	29	0.434	28
Professor	0.158	413	0.162	519
Psychologist	0.370	79	0.336	55
Surgeon	0.290	15	0.249	84
Teacher	0.338	76	0.446	55
<i>DistilBERT — masked</i>				
Accountant	0.413	11	0.425	28
Architect	0.521	14	0.406	47
Attorney	0.223	115	0.220	161
Dentist	0.348	31	0.406	83

(continued on next page)

Profession	Female		Male	
	Faith.	n	Faith.	n
Filmmaker	0.310	17	0.397	34
Journalist	0.311	79	0.235	72
Nurse	0.494	126	0.660	22
Painter	0.476	26	0.368	30
Photographer	0.509	76	0.444	124
Physician	0.183	173	0.201	150
Poet	0.254	28	0.491	26
Professor	0.197	422	0.227	544
Psychologist	0.391	69	0.359	47
Surgeon	0.397	19	0.414	82
Teacher	0.308	75	0.395	48
<i>RoBERTa — unmasked</i>				
Accountant	0.408	12	0.197	29
Architect	0.381	17	0.252	52
Attorney	0.192	108	0.164	156
Dentist	0.289	31	0.265	84
Dietitian	0.549	28	0.533	2
Filmmaker	0.088	20	0.244	37
Journalist	0.259	84	0.309	68
Nurse	0.163	130	0.534	16
Painter	0.310	28	0.367	29
Photographer	0.291	72	0.278	131
Physician	0.280	174	0.184	154
Poet	0.139	28	0.108	28
Professor	0.155	411	0.168	533
Psychologist	0.234	77	0.227	52
Surgeon	0.384	15	0.226	86
Teacher	0.249	85	0.307	46
<i>RoBERTa — masked</i>				
Accountant	0.382	14	0.215	30
Architect	0.405	18	0.221	53
Attorney	0.190	109	0.134	153
Dentist	0.241	31	0.223	85
Filmmaker	0.177	19	0.194	36
Journalist	0.313	77	0.324	74
Nurse	0.172	129	0.374	18
Painter	0.363	28	0.324	32
Photographer	0.418	74	0.302	124
Physician	0.218	170	0.181	154
Poet	0.142	28	0.159	27
Professor	0.172	410	0.191	525
Psychologist	0.274	80	0.221	55
Surgeon	0.247	20	0.214	83
Teacher	0.225	79	0.354	52

Professions where only one gender reaches $n \geq 10$ (comedian|M, composer|M, model|F, software engineer|M) are omitted; their overall faithfulness appears in Appendix E. Dietitian|M ($n=3$ in DU; $n=2$ in RU; $n < 10$ in masked conditions) is included for reference but should be interpreted with caution.

The nurse|M asymmetry remains the largest gender-conditioned faithfulness gap in both unmasked conditions: DistilBERT 0.756 vs. 0.354, RoBERTa 0.534 vs. 0.163. A second notable asymmetry appears for *poet* under DistilBERT unmasked (0.434 vs. 0.200), suggesting that male poet predictions are similarly concentrated on a small decisive feature set—though it does not persist under masking (0.491 vs. 0.254) or under RoBERTa (0.108 vs.

0.139). Under masking, the nurse gap narrows substantially (DistilBERT: 0.660 vs. 0.494; RoBERTa: 0.374 vs. 0.172) but does not disappear, consistent with implicit proxy signals persisting after explicit gender-token removal.

References

- Giuseppe Attanasio, Debora Nozza, Dirk Hovy, and Elena Baralis. Entropy-based attention regularization frees unintended bias mitigation from lists. In *Findings of the Association for Computational Linguistics: ACL 2022*, pages 1105–1119, 2022.
- Stella Biderman, Hailey Schoelkopf, Quentin Gregory Anthony, Herbie Bradley, Kyle O’Brien, Eric Hallahan, Mohammad Aflah Khan, Shivanshu Purohit, USVSN Sai Prashanth, Edward Raff, et al. Pythia: A suite for analyzing large language models across training and scaling. In *Proceedings of the 40th International Conference on Machine Learning*, pages 2397–2430, 2023.
- Rishi Bommasani, Drew A Hudson, Ehsan Adeli, et al. On the opportunities and risks of foundation models. *arXiv preprint arXiv:2108.07258*, 2021.
- Maria De-Arteaga, Alexey Romanov, Hanna Wallach, Jennifer Chayes, Christian Borgs, Alexandra Chouldechova, Sahin Geyik, Krishnaram Kenthapadi, and Adam Tauman Kalai. Bias in bios: A case study of semantic representation bias in a high-stakes setting. In *Proceedings of the ACM Conference on Fairness, Accountability, and Transparency*, pages 120–128, 2019.
- Cynthia Dwork, Moritz Hardt, Toniann Pitassi, Omer Reingold, and Richard Zemel. Fairness through awareness. In *Proceedings of the 3rd Innovations in Theoretical Computer Science Conference*, pages 214–226, 2012.
- Hila Gonen and Yoav Goldberg. Lipstick on a pig: Debiasing methods cover up systematic gender biases in word embeddings but do not remove them. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics*, pages 609–614, 2019.
- Moritz Hardt, Eric Price, and Nati Srebro. Equality of opportunity in supervised learning. In *Advances in Neural Information Processing Systems*, volume 29, 2016.
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. Lora: Low-rank adaptation of large language models. In *Proceedings of the International Conference on Learning Representations*, 2022.
- Yinhan Liu, Mili Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. Roberta: A robustly optimized bert pretraining approach. *arXiv preprint arXiv:1907.11692*, 2019.
- Kaiji Lu, Piotr Mardziel, Fangjing Wu, Preethi Amancharla, and Anupam Datta. Gender bias in neural natural language processing. In *Logic, Language, Information, and Computation*, pages 189–202, 2020.
- Scott M Lundberg and Su-In Lee. A unified approach to interpreting model predictions. In *Advances in Neural Information Processing Systems*, volume 30, 2017.
- Marco Tulio Ribeiro, Sameer Singh, and Carlos Guestrin. ”why should i trust you?”: Explaining the predictions of any classifier. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 1135–1144, 2016.
- Alexey Romanov, Maria De-Arteaga, Hanna Wallach, Jennifer Chayes, Christian Borgs, Alexandra Chouldechova, Sahin Geyik, Krishnaram Kenthapadi, Anna Rumshisky, and Adam Tauman Kalai. What’s in a name? reducing bias in bios without access to protected attributes. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics*, pages 4195–4207, 2019.

Victor Sanh, Lysandre Debut, Julien Chaumond, and Thomas Wolf. Distilbert, a distilled version of bert: smaller, faster, cheaper and lighter. *arXiv preprint arXiv:1910.01108*, 2019.

Mukund Sundararajan, Ankur Taly, and Qiqi Yan. Axiomatic attribution for deep networks. In *Proceedings of the 34th International Conference on Machine Learning*, pages 3319–3328, 2017.